



Project 1: 2 days left

# Prevention-Based Cyber Security: Reference Monitor

CS 459/559: Science of Cyber Security  
12<sup>th</sup> Lecture

**Instructor:**

Guanhua Yan


# Agenda

- ~~Quiz 1: September 29 (closed book)~~
- Project 1 (offense): October 10
- Project 2 (defense): December 5
- Presentations: 11/17, 11/19, 11/24, 12/1, 12/3
- Final report: December 15



# Goal of reference monitor

- Idea: enforce **security policy** at the gateway



# Outline

- **What is reference monitor**
- **Power of reference monitor**
- **Examples of reference monitor**
  - Hardware-based protection
  - Software fault isolation
  - Control flow integrity




# Reference Monitors



## Reference Monitor

- Observe the execution of a program and halt the program if it's going to violate the security policy.





## Common Examples of RM

- Operating system monitors users applications
  - Monitor system calls by user apps
  - Kernel vs user mode
  - Hardware based
- Software-based: Interpreters, language virtual machines, software-based fault isolation
- Firewalls
- ...
- Claim: majority of today's security enforcement mechanisms are instances of reference monitors



## Requirements for a Monitor

- Must have (reliable) access to information about what the program is about to do
  - e.g., what syscall is it about to execute?
- Must have the ability to “stop” the program
  - can’t stop a program running on another machine that you don’t own
  - stopping isn’t necessary; transitioning to a “good” state may be sufficient
- Must protect the monitor’s state and code from tampering
  - key reason why a kernel’s data structures and code aren’t accessible by user code
- In practice, must have low overhead



## Requirements for a Monitor

- Must have (reliable) access to information about what the program is about to do
  - e.g., what syscall is it about to execute?
- Must have the ability to “stop” the program
  - can’t stop a program running on another machine that you don’t own
  - stopping isn’t necessary; transitioning to a “good” state may be sufficient
- Must protect the monitor’s state and code from tampering
  - key reason why a kernel’s data structures and code aren’t accessible by user code
- In practice, must have low overhead



## Requirements for a Monitor

- Must have (reliable) access to information about what the program is about to do
  - e.g., what syscall is it about to execute?
- Must have the ability to “stop” the program
  - can’t stop a program running on another machine that you don’t own
  - stopping isn’t necessary; transitioning to a “good” state may be sufficient
- Must protect the monitor’s state and code from tampering
  - key reason why a kernel’s data structures and code aren’t accessible by user code
- In practice, must have low overhead



## Requirements for a Monitor

- Must have (reliable) access to information about what the program is about to do
  - e.g., what syscall is it about to execute?
- Must have the ability to “stop” the program
  - can’t stop a program running on another machine that you don’t own
  - stopping isn’t necessary; transitioning to a “good” state may be sufficient
- Must protect the monitor’s state and code from tampering
  - key reason why a kernel’s data structures and code aren’t accessible by user code
- In practice, must have low overhead



# What Policies Can be Enforced?

- Some liberal assumptions:

- Monitor can have infinite state
- Monitor can have access to entire history of computation
- But monitor can't guess the future – the decision to determine whether to halt a program must be computable

- Under these assumptions:

- There is a nice class of policies that reference monitors can enforce: **safety properties**
- There are desirable policies that no reference monitor can enforce precisely



## What Policies Can be Enforced?

- Some liberal assumptions:
  - Monitor can have infinite state
  - Monitor can have access to entire history of computation
  - But monitor can't guess the future – the decision to determine whether to halt a program must be computable
- Under these assumptions:
  - There is a nice class of policies that reference monitors can enforce: **safety properties**
  - There are desirable policies that no reference monitor can enforce precisely



## What Policies Can be Enforced?

- Some liberal assumptions:
  - Monitor can have infinite state
  - Monitor can have access to entire history of computation
  - But monitor can't guess the future – the decision to determine whether to halt a program must be computable
- Under these assumptions:
  - There is a nice class of policies that reference monitors can enforce: **safety properties**
  - There are desirable policies that no reference monitor can enforce precisely



## What Policies Can be Enforced?

- Some liberal assumptions:
  - Monitor can have infinite state
  - Monitor can have access to entire history of computation
  - But monitor can't guess the future – the decision to determine whether to halt a program must be computable
- Under these assumptions:
  - There is a nice class of policies that reference monitors can enforce: **safety properties**
  - There are desirable policies that no reference monitor can enforce precisely



# **Analysis of the Power and Limitations of Execution Monitoring**


“Enforceable Security Policies” by Fred Schneider





# Execution Traces


- System behavior  $\sigma$ : a finite or infinite execution trace of system events
  - $\sigma = e_0 e_1 e_2 e_3 \dots e_i \dots$ , where  $e_i$  is a system event
- Example: a trace of memory operations (reads and writes)
  - Events: `read(addr)`; `write(addr, v)`
- Example: a trace of system calls
  - System-call events: `open(...)`; `read(...)`; `close(...)`; `gettimeofday(...)`; `fork(...)`; ...
- Example: a system for access control
  - `Oper(p, o, r)`: Principal p invoked an operation involving object o and requiring right r to that object
  - `AddP(p, p')`: Principal p invoked an operation to create a principal named p'
  - ...





# Execution Traces


- System behavior  $\sigma$ : a finite or infinite execution trace of system events
  - $\sigma = e_0 e_1 e_2 e_3 \dots e_i \dots$ , where  $e_i$  is a system event
- Example: a trace of memory operations (reads and writes)
  - Events: `read(addr)`; `write(addr, v)`
- Example: a trace of system calls
  - System-call events: `open(...)`; `read(...)`; `close(...)`; `gettimeofday(...)`; `fork(...)`; ...
- Example: a system for access control
  - `Oper(p, o, r)`: Principal p invoked an operation involving object o and requiring right r to that object
  - `AddP(p, p')`: Principal p invoked an operation to create a principal named p'
  - ...





# Execution Traces

- System behavior  $\sigma$ : a finite or infinite execution trace of system events
  - $\sigma = e_0 e_1 e_2 e_3 \dots e_i \dots$ , where  $e_i$  is a system event
- Example: a trace of memory operations (reads and writes)
  - Events: `read(addr)`; `write(addr, v)`
- Example: a trace of system calls
  - System-call events: `open(...)`; `read(...)`; `close(...)`; `gettimeofday(...)`; `fork(...)`; ...
- Example: a system for access control
  - `Oper(p, o, r)`: Principal p invoked an operation involving object o and requiring right r to that object
  - `AddP(p, p')`: Principal p invoked an operation to create a principal named p'
  - ...






# Execution Traces

- System behavior  $\sigma$ : a finite or infinite execution trace of system events
  - $\sigma = e_0 e_1 e_2 e_3 \dots e_i \dots$ , where  $e_i$  is a system event
- Example: a trace of memory operations (reads and writes)
  - Events: `read(addr)`; `write(addr, v)`
- Example: a trace of system calls
  - System-call events: `open(...)`; `read(...)`; `close(...)`; `gettimeofday(...)`; `fork(...)`; ...
- Example: a system for access control
  - `Oper(p, o, r)`: Principal p invoked an operation involving object o and requiring right r to that object
  - `AddP(p, p')`: Principal p invoked an operation to create a principal named p'

...





## Modeling a System

- A **system** modeled as a set of execution traces (its behaviors)
  - $S = \{\sigma_1, \sigma_2, \dots, \sigma_k, \dots\}$
  - Each trace corresponds to the execution for a possible input

- For example
  - Sets of traces of reads and writes
  - Sets of traces of system calls



# Definition of Security Policy

- A **security policy**  $P(S)$ : a logical predicate on sets of execution traces.
  - A target system  $S$  satisfies security policy  $P$  if and only if  $P(S)$  holds.
- For example
  - A program cannot write to addresses outside of  $[0, 1000]$ 
    - $P(S) = \forall \sigma \in S. \forall e_i \in \sigma. e_i = \text{write}(addr, v) \rightarrow addr \in [0, 1000]$
  - A program cannot send a network packet after reading from a local file
    - $P(S) = \forall \sigma \in S. \forall e_i \in \sigma. e_i = \text{fileRead}(\dots) \rightarrow \forall k > i. e_k \neq \text{networkSend}(\dots)$



## Constraint on Monitors: Property

- Can a reference monitor see more than one trace at a time?
  - A reference monitor only sees one execution trace of a program
- So we can only enforce policies  $P$  s.t.:
  - (1)  $P(S) = \forall \sigma \in S. P(\sigma)$
  - where  $P$  is a predicate on individual traces
- A security policy is a **property** if its predicate specifies whether an individual trace is legal
  - The membership is determined solely by the trace and not by the other traces

## Enforceability



# What is a Non-Property?

- A policy that may depend on multiple execution traces
- Information flow policies
  - Sensitive information should not flow to unauthorized person explicitly or **implicitly**
  - Example: a system protected by passwords
    - Suppose the password checking time correlates closely to the length of the prefix that matches the true password
      - Timing channel
    - To rule this out, a policy should say: no matter what the input is, the password checking time should be the same in **all traces**
    - Not a property



## More Constraints on Monitors

Shouldn't be able to "see" the future.

- Assumption: must make decisions in finite time.
- Suppose  $\mathcal{P}(\sigma)$  is false, then it must be rejected at some finite time  $i$ ; that is,  $\mathcal{P}(\sigma[..i])$  is false

$$(2) \forall \sigma. \neg \mathcal{P}(\sigma) \rightarrow (\exists i. \neg \mathcal{P}(\sigma[..i]))$$

Once a trace has been rejected by the monitor, then any further events from the system cannot make the monitor to revoke that decision

$$(3) \forall \sigma. \neg \mathcal{P}(\sigma) \rightarrow (\forall \sigma'. \neg \mathcal{P}(\sigma\sigma'))$$

## Prefix Closure



## More Constraints on Monitors

Shouldn't be able to "see" the future.

- Assumption: must make decisions in finite time.
- Suppose  $\mathcal{P}(\sigma)$  is false, then it must be rejected at some finite time  $i$ ; that is,  $\mathcal{P}(\sigma[..i])$  is false

$$(2) \forall \sigma. \neg \mathcal{P}(\sigma) \rightarrow (\exists i. \neg \mathcal{P}(\sigma[..i]))$$

Once a trace has been rejected by the monitor, then any further events from the system cannot make the monitor to revoke that decision

$$(3) \forall \sigma. \neg \mathcal{P}(\sigma) \rightarrow (\forall \sigma'. \neg \mathcal{P}(\sigma\sigma'))$$

# Finite refutability



## Reference Monitors Enforce Safety Properties

A predicate  $P$  on sets of sequences s.t.

- (1)  $P(S) = \forall \sigma \in S. P(\sigma)$
- (2)  $\forall \sigma. \neg P(\sigma) \rightarrow (\exists i. \neg P(\sigma[..i]))$
- (3)  $\forall \sigma. \neg P(\sigma) \rightarrow (\forall \sigma'. \neg P(\sigma\sigma'))$

is a *safety property*: “no bad thing will happen.”

Conclusion: a reference monitor can't enforce a policy  $P$  unless it's a safety property.




## Safety and Liveness Properties [Alpern & Schneider 85,87]

- Safety: Some “bad thing” doesn’t happen.
  - Proscribes traces that contain some “bad” prefix
  - Example: the program won’t read memory outside of range [0,1000]
- Liveness: Some “good thing” does happen
  - Example: program will terminate
  - Example: program will eventually release the lock
- Theorem: Every security property can be decomposed into a safety property and a liveness property



## Classification of Policies

- “Enforceable Security Policies” [Schneider 00]





## Policies Enforceable by Reference Monitors

- Reference monitor **can** enforce any safety property
  - Intuitively, the monitor can inspect the history of computation and prevent bad things from happening
- Reference monitor **cannot** enforce liveness properties
  - The monitor cannot predict the future of computation
- Reference monitor **cannot** enforce non-properties
  - The monitor cannot inspect multiple traces simultaneously




# Safety Is Nice

Safety has its benefits:

- They compose: if P and Q are safety properties, then P & Q is a safety property (just the intersection of allowed traces.)
- Safety properties can approximate liveness by setting limits. e.g., we can determine that a program terminates within k steps.
- We can also approximate many other security policies (e.g., info. flow) by simply choosing a stronger safety property.

# **Reference Monitor Example 1: Hardware-Based Protection**





# Operating Systems circa '75

Simple Model: system is a collection of running processes and files.

- processes perform actions on behalf of a user.
  - open, read, write files
  - read, write, execute memory, etc.
- files have access control lists dictating which users can read/write/execute/etc. the file.

(Some) High-Level Policy Goals:

- Integrity: one user's processes shouldn't be able to corrupt the code, data, or files of another user.
- Availability: processes should eventually gain access to resources such as the CPU or disk.
- Confidentiality? Access control?




## What Can go Wrong?

- read/write/execute or change ACL of a file for which process doesn't have proper access.
  - check file access against ACL
- process writes into memory of another process
  - isolate memory of each process (& the OS!)
- process pretends it is the OS and execute its code
  - maintain process ID and keep certain operations privileged --- need some way to transition.
- process never gives up the CPU
  - force process to yield in some finite time
- process uses up all the memory or disk
  - enforce quotas



# Hardware-Based Protection

- Based on virtual memory
  - Protect one process from reading/writing to other processes' memory locations
  - Memory safety (fault isolation) at the process level





# Virtual Memory

- Each process assumes a virtual address space
- A page table translates virtual pages to physical pages
  - Dedicated hardware for acceleration: TLB

virtual page	physical page	permission
0	0xABCD0000	rw
1	0xA0700000	r
:		

A **translation lookaside buffer (TLB)** is a memory cache that stores the recent translations of virtual memory to physical memory. It is used to reduce the time taken to access a user memory location.



## Hardware Privilege Modes


- Generally two: kernel mode & user mode  
(x86 has more)
- At any time, CPU is in some mode
- Dangerous (“privileged”) instructions usable  
only in the kernel mode
  - E.g., direct access to any physical memory  
location, or privilege-mode change
  - Violations trap to a fixed address in the kernel



## Moving Between Privilege Modes

- Downgrade privilege
  - Usually have a simple instruction for this
- Upgrade privilege
  - Usually have a special “system call” instruction (like function call), but
    - Go to a fixed address or through a fixed jumptable
    - Change into privileged (kernel) mode

## Steps in a System Call





# Pros and Cons of Hardware-Based Protection

## ○ Pros

- Low overhead; built into the hardware
- Transparent for applications

## ○ Cons

- Inter-process communication is cumbersome and slow
  - Signals, remote procedural calls (RPC), pipes, sockets, shared memory, ...
  - Involves context switch (changes to a new page table)
- Granularity of protection is per-process
  - Doesn't protect a buffer of size, say, 100 bytes
  - Buffer overflow attacks apply

# **Reference Monitor Example 2: Software Fault Isolation**

# SFI Example: Google Chrome Sandbox

## A new approach to browser security: the Google Chrome Sandbox

Thursday, October 2, 2008

Building a secure browser is a top priority for the Chromium team; it's why we spend a lot of time and effort keeping our code secure. But as you can imagine, code perfection is something almost impossible to achieve for a project of this size and complexity. To make things worse, a browser spends most of its time handling and executing untrusted and potentially malicious input data. In the event that something goes wrong, the team has developed a sandbox to help thwart any exploit in two of the most popular vectors of attack against browsers: HTML Rendering and JavaScript execution.

In a nutshell, a sandbox is security mechanism used to run an application in a restricted environment. If an attacker is able to exploit the browser in a way that lets him run arbitrary code on the machine, the sandbox would help prevent this code from causing damage to the system. The sandbox would also help prevent this exploit from modifying and even reading your files or any information on the system.

<https://blog.chromium.org/2008/10/new-approach-to-browser-security-google.html>

# SFI Goals

## ■ **Confine faults inside distrusted extensions**

- codec shouldn't compromise media player
- device driver shouldn't compromise kernel
- plugin shouldn't compromise web browser


## ■ **Allow for efficient cross-domain calls**

- numerous calls between media player and codec
- numerous calls between device driver and kernel

# Software Fault Isolation


- ◆ SFI approach:

- Partition process memory into segments




- Locate unsafe instructions: jmp, load, store
    - At compile time, add guards before unsafe instructions
    - When loading code, ensure all guard are present

# Main Idea



# Scheme 1: Segment Matching

- Check every memory access for matching segment id
- Assume dedicated registers **segment register (sr)** and **data register (dr)**
  - not available to the program




# Safety

- Segment matching code must always be run to ensure safety.
- Dedicated registers must not be writeable by module.

# Scheme 2: Sandboxing

- Force top bits to match seg id and continue
- No comparison is made



# Segment Matching vs. Sandboxing

## Segment Matching

- more instructions
- can pinpoint exact point of fault where segment id doesn't match

## Sandboxing

- fewer instructions
- just ensures memory access stays in region (crash is ok)

# Reference Monitor Example 3: Control Flow Integrity

# Microsoft Control Flow Guard

## What is Control Flow Guard?

Control Flow Guard (CFG) is a highly-optimized platform security feature that was created to combat memory corruption vulnerabilities. By placing tight restrictions on where an application can execute code from, it makes it much harder for exploits to execute arbitrary code through vulnerabilities such as buffer overflows. CFG extends previous exploit mitigation technologies such as [/GS](#), [DEP](#), and [ASLR](#).

This feature is available in Microsoft Visual Studio 2015, and runs on "CFG-Aware" versions of Windows—the x86 and x64 releases for Desktop and Server of Windows 10 and Windows 8.1 Update (KB3000850).

# LLVM Clang

The screenshot shows a web browser window displaying the LLVM Clang documentation for Control Flow Integrity. The URL in the address bar is `clang.llvm.org`. The page title is "Clang 7 documentation" and the specific section is "CONTROL FLOW INTEGRITY". Below the title, there is a breadcrumb navigation: "« Sanitizer special case list :: Contents :: Control Flow Integrity Design Documentation »". The main content area is titled "Control Flow Integrity" and contains a table of contents (TOC) for various CFI schemes:

- [Introduction](#)
- [Available schemes](#)
- [Trapping and Diagnostics](#)
- [Forward-Edge CFI for Virtual Calls](#)
  - [Performance](#)
- [Bad Cast Checking](#)
- [Non-Virtual Member Function Call Checking](#)
  - [Strictness](#)
- [Indirect Function Call Checking](#)
  - [-fsanitize-cfi\\_icall-generalize-pointers](#)
  - [-fsanitize=cfi\\_icall and -fsanitize=function](#)
- [Blacklist](#)
- [Shared library support](#)
- [Design](#)
- [Publications](#)

**Introduction**

Clang includes an implementation of a number of control flow integrity (CFI) schemes, which are designed to abort the program upon detecting certain forms of undefined behavior that can potentially allow attackers to subvert the program's control flow. These schemes have been optimized for performance, allowing developers to enable them in release builds.


To enable Clang's available CFI schemes, use the flag `-fsanitize=cfi`. You can also enable a subset of available [schemes](#). As currently implemented, all schemes rely on link-time optimization (LTO); so it is required to specify `-fno-lto`, and the linker used must support LTO, for example via the [gold plugin](#).

<https://clang.llvm.org/docs/ControlFlowIntegrity.html>

# Basic Block

**Definition of Basic Block:** A consecutive sequence of instructions / code such that


- the instruction in each position always executes before (dominates) all those in later positions, and
- no outside instruction can execute between two instructions in the sequence



# Basic Block in CFG

**Definition of Basic Block:** A consecutive sequence of instructions / code such that

- control is “straight”  
**(no jump targets except at the beginning,  
no jumps except at the end)**



# CFG Definition

A static ***Control Flow Graph*** is a graph where

- each vertex  $v_i$  is a basic block, and
- there is an edge  $(v_i, v_j)$  if there ***may*** be a transfer of control from block  $v_i$  to block  $v_j$ .

Historically, the scope of a “CFG” is limited to a function or procedure, i.e., *intra*-procedural.

# Build CFG


```

bool lt(int x, int y) {
    return x < y;
}

bool gt(int x, int y) {
    return x > y;
}

sort2(int a[], int b[], int len)
{
    sort( a, len, lt );
    sort( b, len, gt );
}

```



# Instrument Binary


```

bool lt(int x, int y) {
    return x < y;
}

bool gt(int x, int y) {
    return x > y;
}

sort2(int a[], int b[], int len)
{
    sort( a, len, lt );
    sort( b, len, gt );
}

```



- Insert a unique number at each destination
- Two destinations are equivalent if CFG contains edges to each from the same source

# Example of instrumentation

Function Call		Function Return	
Opcode bytes	Instructions	Opcode bytes	Instructions
FF 53 08	call [ebx+8] ; call fptr	C2 10 00	ret 10h ; return
are instrumented using prefetchnta destination IDs, to become			
8B 43 08	mov eax, [ebx+8] ; load fptr	8B 0C 24	mov ecx, [esp] ; load ret
3E 81 78 04 78 56 34 12	cmp [eax+4], 12345678h ; comp w/ID	83 C4 14	add esp, 14h ; pop 20
75 13	jne error_label ; if != fail	3E 81 79 04	cmp [ecx+4], ; compare
FF D0	call eax ; call fptr	DD CC BB AA	AABBCCDDh ; w/ID
3E 0F 18 05 DD CC BB AA	<b>prefetchnta [AABBCCDDh]</b> ; label ID	75 13	jne error_label ; if!=fail
		FF E1	jmp ecx ; jump ret

Abuse an x86 assembly instruction to insert "12345678" tag into the binary

Jump to the destination only if the tag is equal to "AABBCCDD"

**Prefetchnta:** prefetch data into non-temporal cache structure and into a location close to the processor

# Verify CFI Instrumentation

- Direct jump targets (e.g. call 0x12345678)
  - are all targets valid according to CFG?
- IDs
  - is there an ID right after every entry point?
  - does any ID appear in the binary by accident?
- ID Checks
  - is there a check before every control transfer?
  - does each check respect the CFG?

easy to implement correctly => trustworthy

# CFI: Preventing Circumvention

## ■ Unique IDs

- Bit patterns chosen as destination IDs must not appear anywhere else in the code memory except ID checks

## ■ Non-writable code

- Program should not modify code memory at runtime
  - What about run-time code generation and self-modification?

## ■ Non-executable data

- Program should not execute data as if it were code

## ■ Enforcement: hardware support + prohibit system calls that change protection state + verification at load-time

# CFI: Security Guarantees

- **Effective against attacks based on illegitimate control-flow transfer**
  - Stack-based buffer overflow, return-to-libc exploits, pointer subterfuge
- **Does not protect against attacks that do not violate the program's original CFG**
  - Incorrect arguments to system calls
  - Substitution of file names
  - Other data-only attacks

*End of Lecture 12*