



Project 3: **39** days left

Detection-Based Cybersecurity: Detection Evasion

CS 459/559: Science of Cyber Security
17th Lecture

Instructor:
Guanhua Yan

Agenda

- ~~Quiz 1: September 29 (closed book)~~
- ~~Project 1 (offense): October 10~~
- Quiz 2: November 12
- Presentations: 11/17, 11/19, 11/24, 12/1, 12/3
- CTF competition: November 26
- Project 2 (defense): December 5
- Final report: December 15



Outline

- Pros and cons of detection techniques
- Two types of intrusion detection systems
- Detection evasion

Pros and cons of detection techniques

Intrusion detection methodology

- Signature-based (misuse-based)
- Anomaly detection
- Specification-based

Misuse Detection

Goal

Detect **known** attacks via *signatures/pattern* or *black lists*

Pros

- ✓ Easy to understand, readily shareable
- ✓ FPs: management likes warm fuzzy feeling

Cons

- ✗ Polymorphism: unable to detect new attacks or variants
- ✗ Accuracy: finding sweetspot between FPs and FNs is *hard*

Example

Snort, regular expression matching

Anomaly Detection

Goal

Flag **deviations** from a known profile of “normal”

Pros

- ✓ Detect wide range of attacks
- ✓ Detect novel attacks

Cons

- ✗ High FP rate
- ✗ Efficacy depends on training data purity

Example

Look at distribution of characters in URLs, learn some are rare

Specification-Based Detection

Goal

Describe what constitutes allowed activity via *policy* or *white list*

Pros

- ✓ Can detect novel attacks
- ✓ Can have low FPs

Cons

- ✗ Expensive: requires significant development
- ✗ Churn: must be kept up to date

Example

Firewall

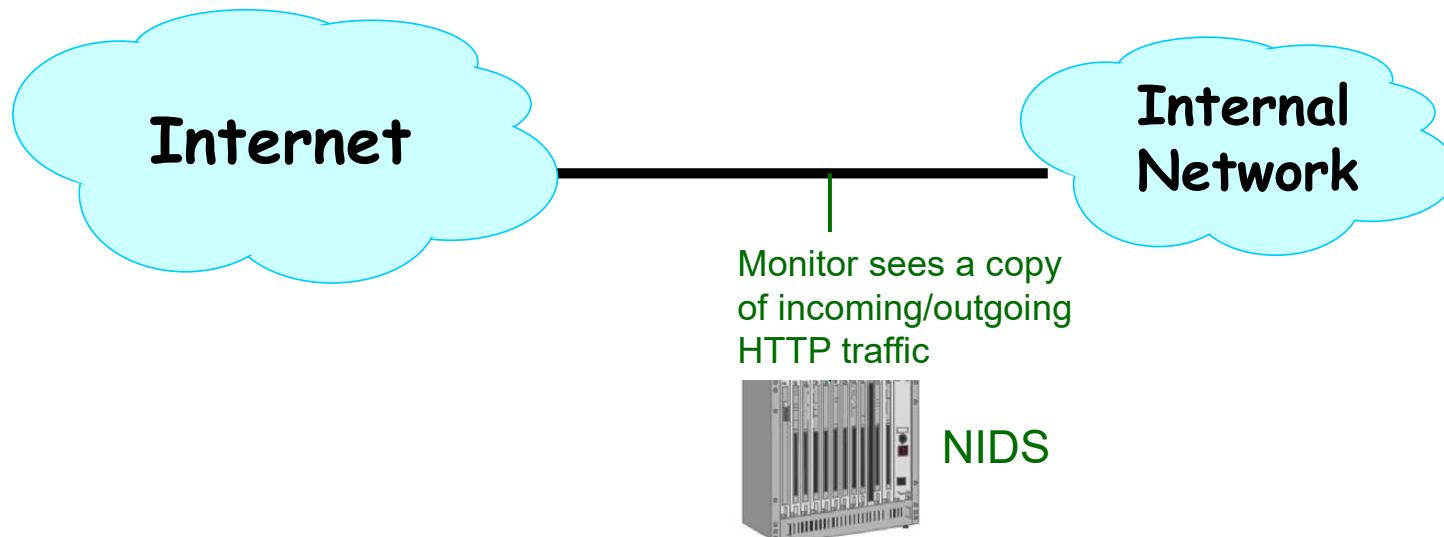
Two types of intrusion detection systems

At least two types of IDS

- NIDS: sits in the network
- HIDS: sits at the end host

Network Intrusion Detection (NIDS)

- Passively monitor network traffic for signs of attack at perimeter of a network
 - Look for certain rules (e.g., /etc/passwd)
 - Flag a warning to an administrator, do not take preemptive action



NIDS rules set

- A set of rules (string matching, regular expression) that identifies an attack
- Example rule:
 - “any flow containing /etc/password should be flagged”
 - “any flow containing attack.exe should be flagged”

What does a NIDS aim to detect?

Examples:

- Port scans: information gathering intended to determine which ports are open for TCP connections
- DoS attacks
- Malware (replicating malicious software)
- DNS cache poisoning
- ARP spoofing

Inside a Modern NIDS

- Deployment **inside** network as well as at border
 - Greater visibility, including **tracking of user identity**
- Full protocol analysis
 - Including extraction of complex embedded objects
 - In some systems, 100s of known protocols
- Signature analysis (also behavioral)
 - Known attacks, malware communication, blacklisted hosts/domains
 - Known malicious payloads
 - Sequences/patterns of activity
- **Shadow execution** (e.g., Flash, PDF programs)
- Extensive logging (in support of **forensics**)
- Auto-update of signatures, blacklists

Inside a Modern HIDS

- *Sandbox execution*
 - Run selected executables in constrained/monitored environment
 - Analyze:
 - System calls
 - Changes to files / registry
 - Self-modifying code (*polymorphism/metamorphism*)
- File scanning
 - Look for malware that installs itself on disk
- Memory scanning
 - Look for malware that **never appears on disk**
- Runtime analysis
 - Apply heuristics/signatures to execution behavior

NIDS vs. HIDS

- NIDS benefits:
 - Can **cover a lot of systems** with single deployment
 - Much simpler management
 - Easy to “bolt on” / **no need to touch end systems**
 - Doesn’t consume production resources on end systems
 - Harder for an attacker to subvert / less to trust
- HIDS benefits:
 - Can have **direct access to semantics** of activity
 - Better positioned to block (prevent) attacks
 - Harder to evade
 - Can protect against non-network threats
 - **Visibility** into encrypted activity
 - Performance scales much more readily (no chokepoint)
 - No issues with “dropped” packets

Detection evasion

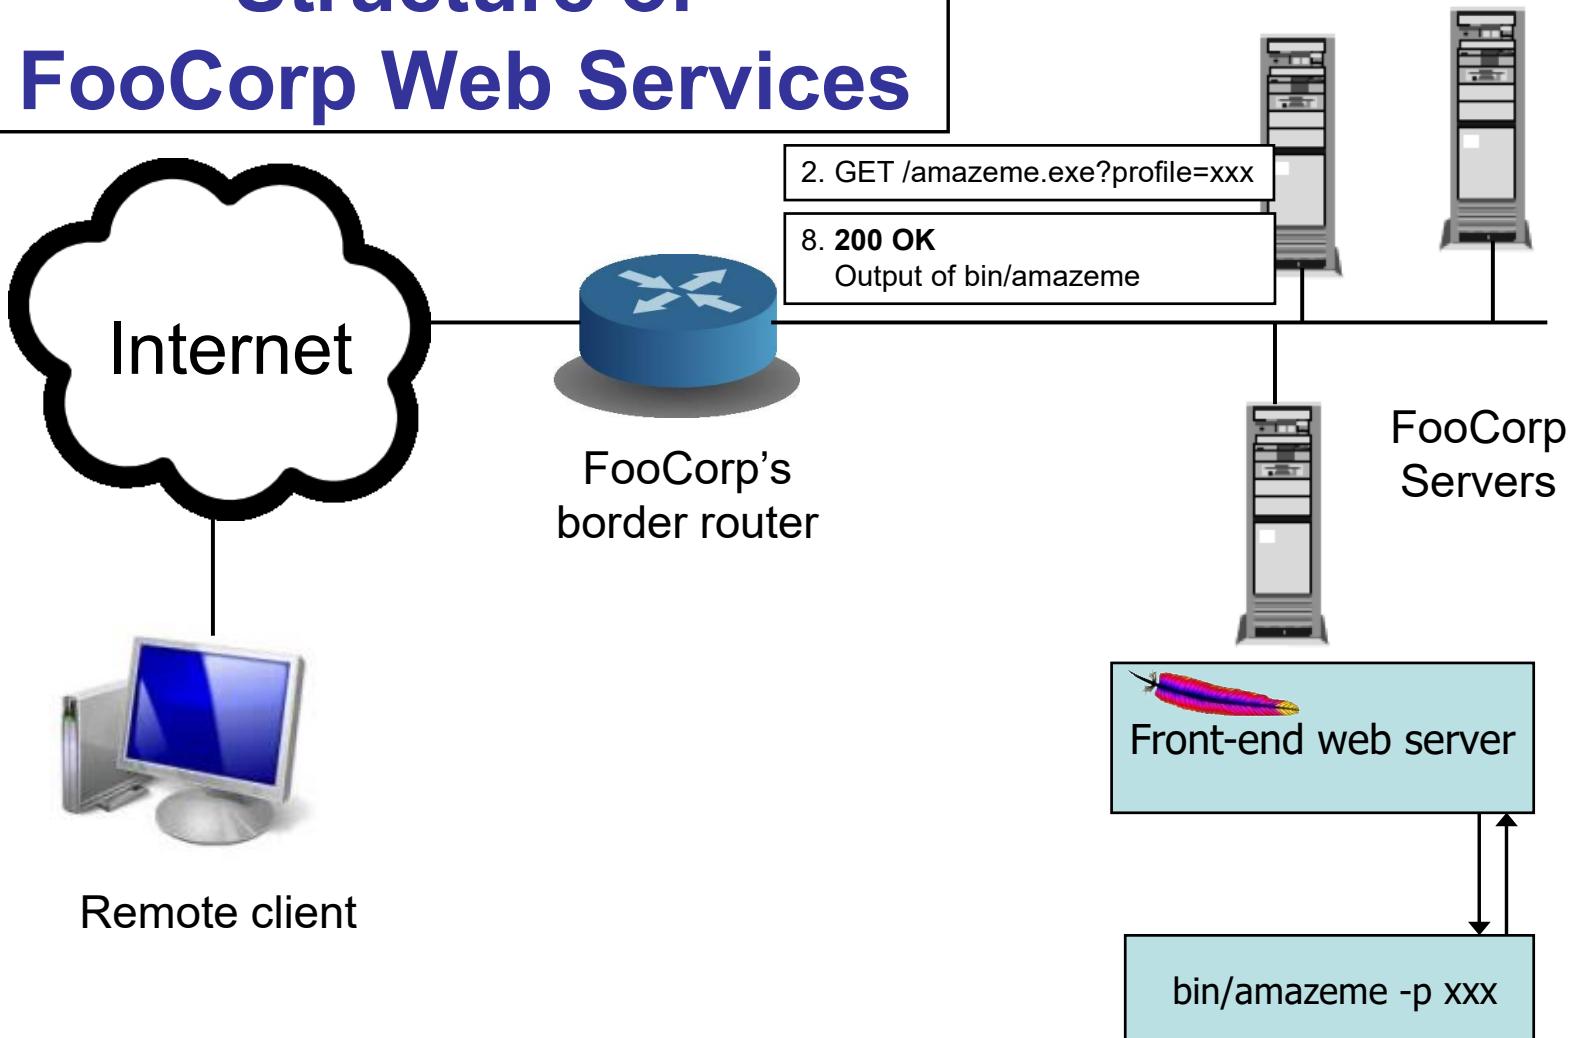
Evasion

Evasion attacks can arise when you have
“double parsing”

- *Inconsistency* – interpreted differently
- *Ambiguity* – information needed to interpret
is missing

Or when you attack the IDS

Structure of FooCorp Web Services



Evasion Attacks (High-Level View)

- Some evasions reflect **incomplete analysis**
 - In our FooCorp example, hex escapes or “..///.//..” alias
 - In principle, can deal with these with implementation care (make sure we **fully understand the spec**)
- Some are due to **imperfect observability**
 - For instance, if what NIDS sees doesn’t exactly match what arrives at the destination
- Some are due to **attacking the IDS itself**

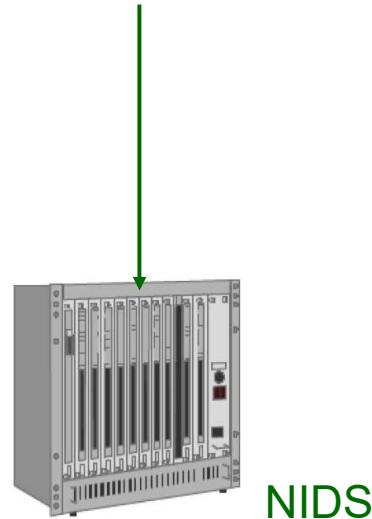
Evasion

- What should NIDS do if it sees a RST packet?



- (a) Assume RST will be received
- (b) Assume RST won't be received

Safer to consider both possibilities



Evasion

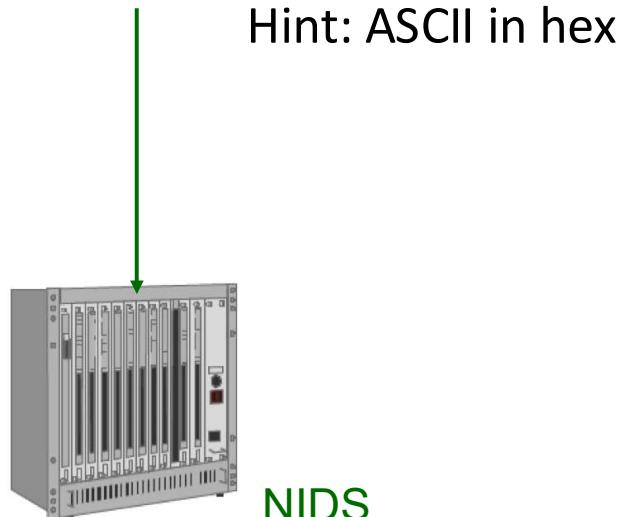
- What should NIDS do if it sees this?

/%65%74%63/%70%61%73%73%77%64

Hint: ASCII in hex

- (a) Alert – it's an attack
- (b) No alert – it's all good
- (c) Other (please specify)

This can be /etc/passwd depending
on what protocol parses this,
ideally it would realize it is an
attack and alert

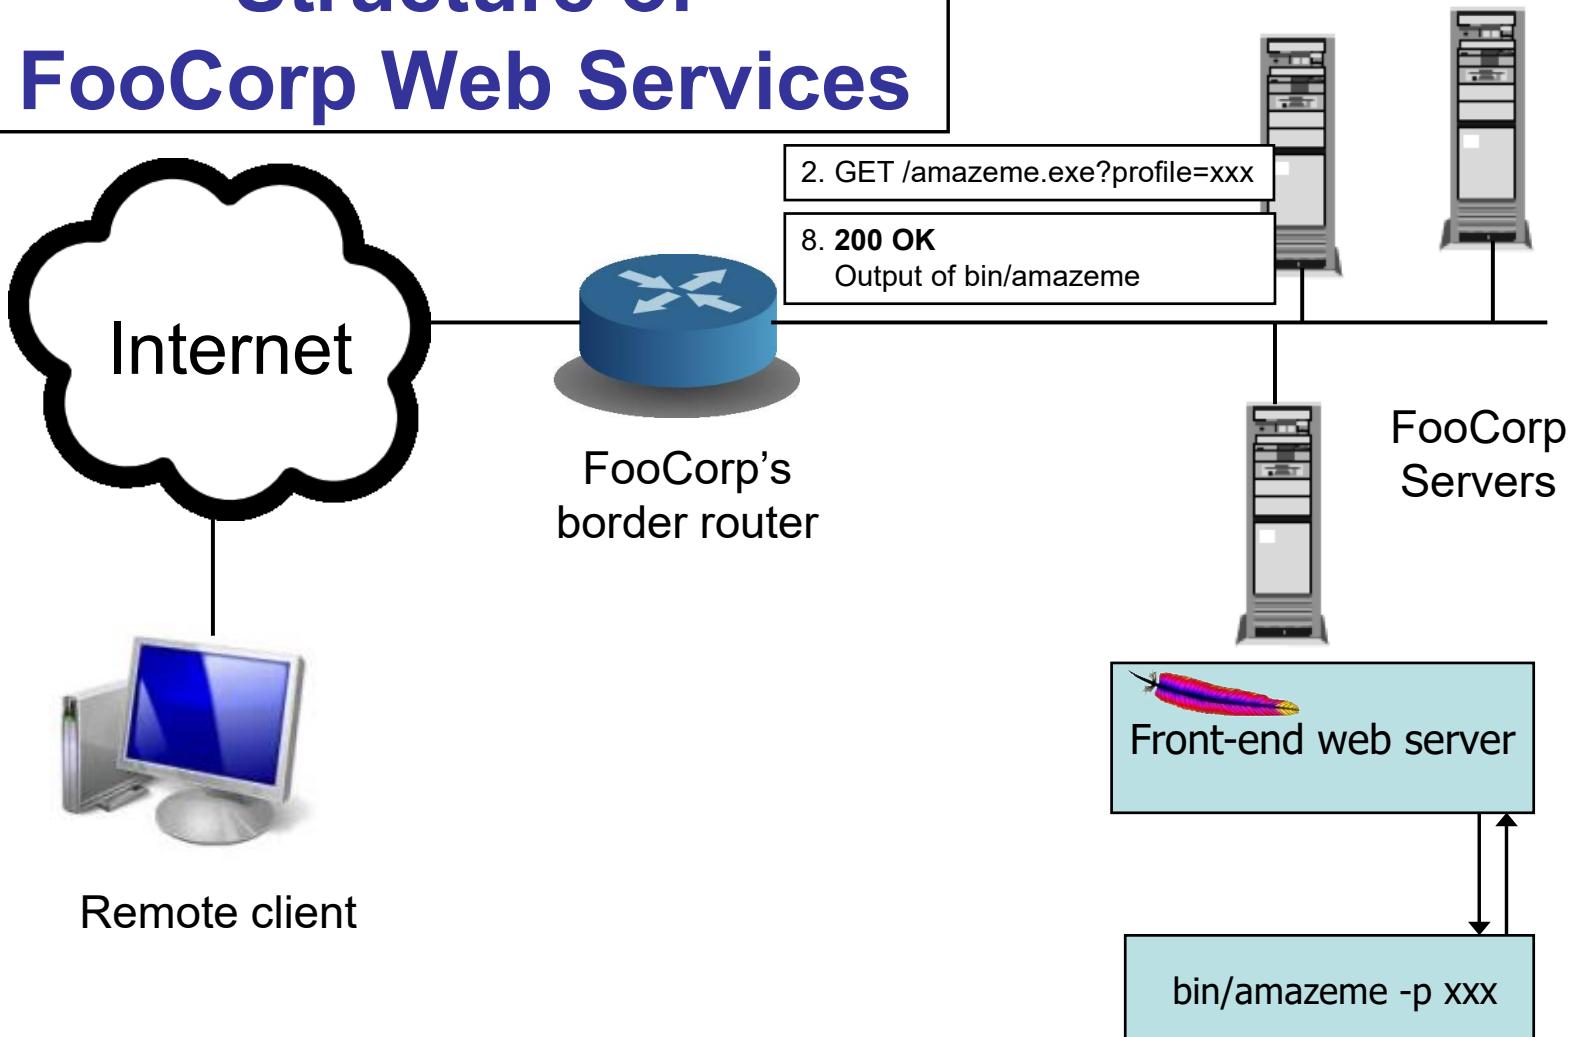


Evasion

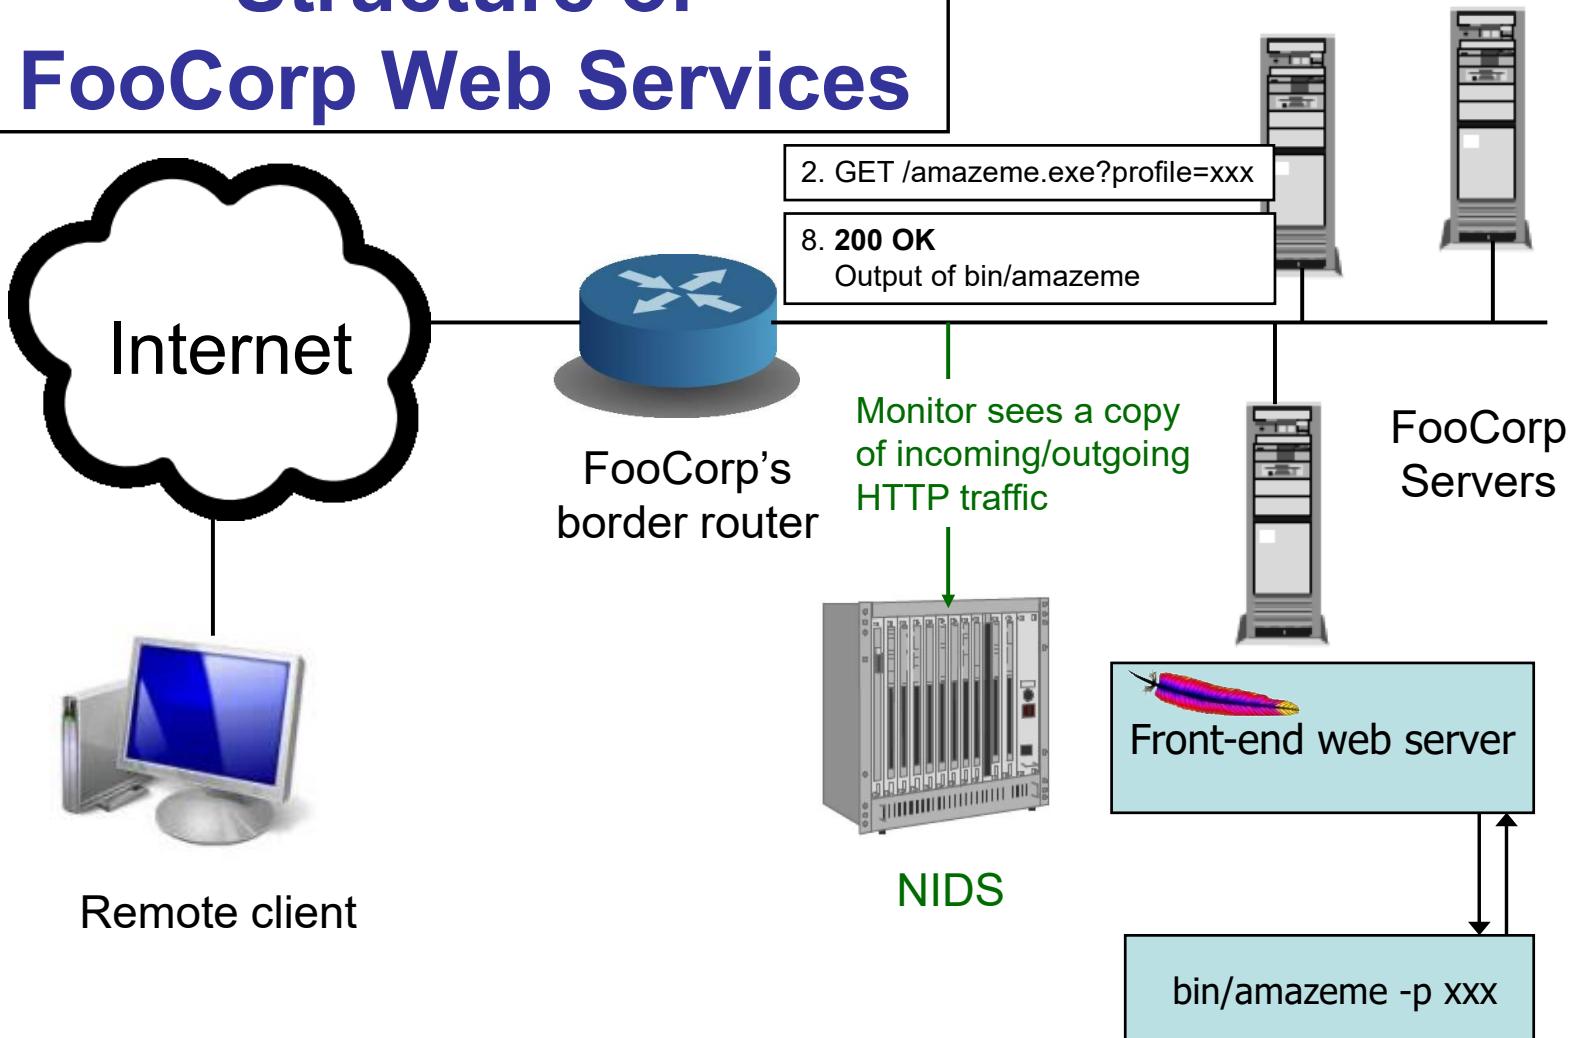
How can you mount a DoS on the IDS?

- Send so many attacks that matches rules to the IDS making the IDS log so much data that it becomes slow or runs out of resources
- Or fake new connections so the IDS creates new state

Structure of FooCorp Web Services



Structure of FooCorp Web Services



Network Intrusion Detection

- Approach #1: look at the network traffic
 - (a “NIDS”: rhymes with “kids”)
 - Scan HTTP requests
 - Look for “/etc/passwd” and/or “..../”

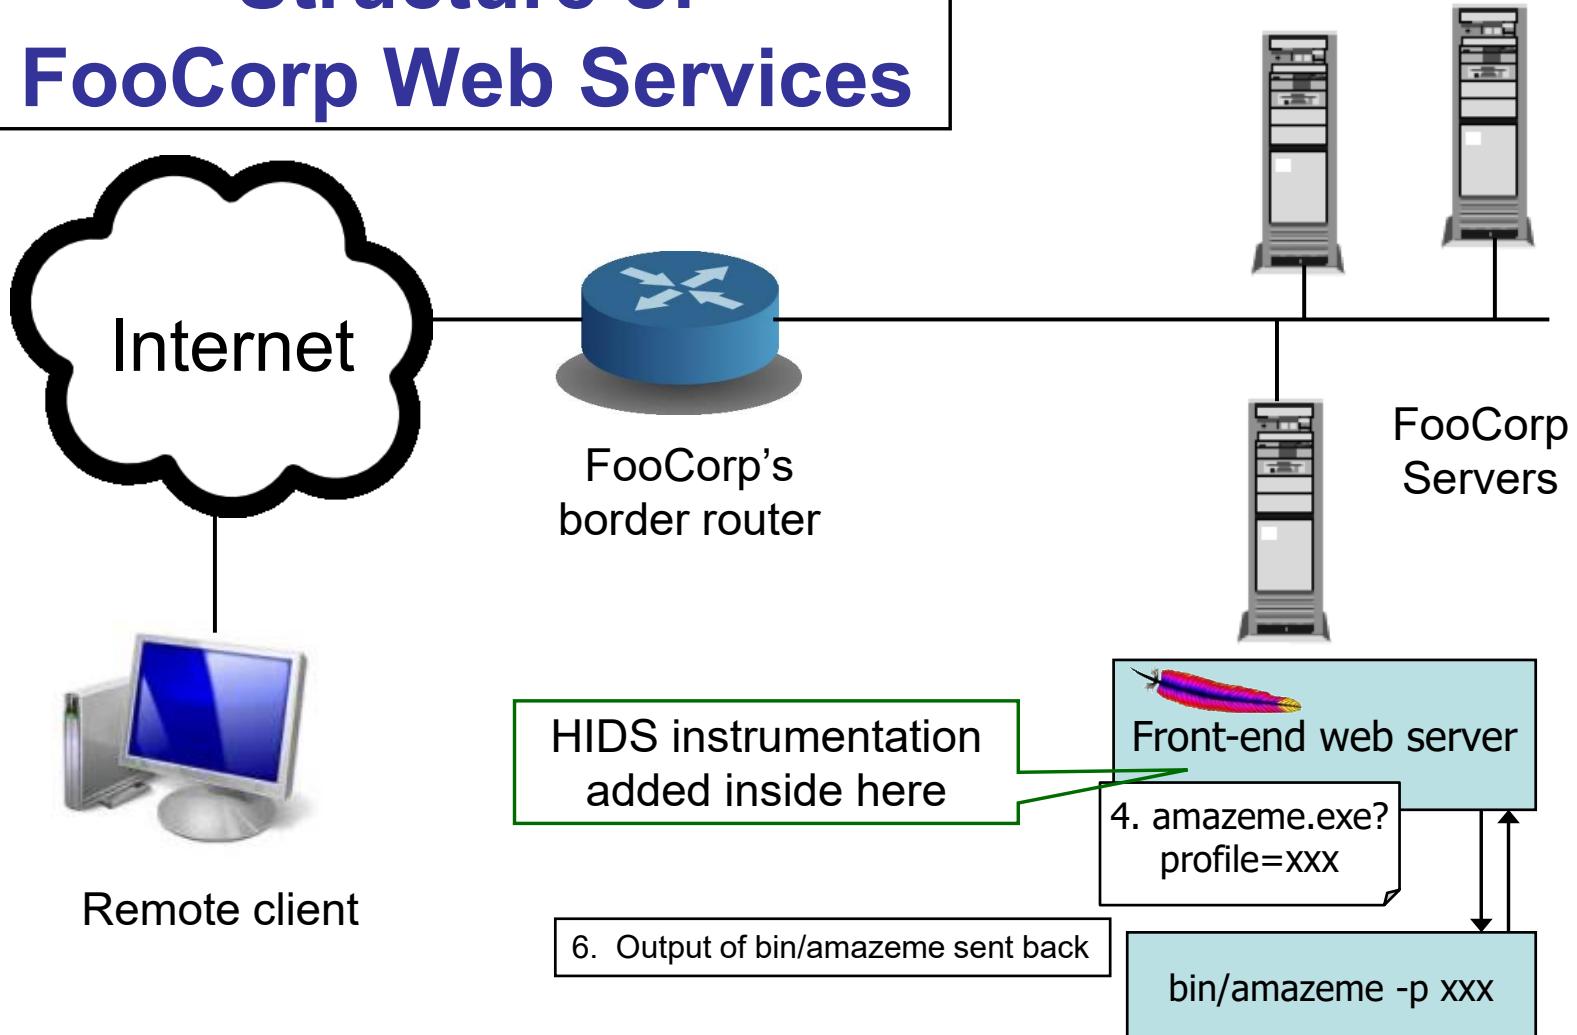
Network Intrusion Detection

- Approach #1: look at the network traffic
 - (a “NIDS”: rhymes with “kids”)
 - Scan HTTP requests
 - Look for “/etc/passwd” and/or “..../”
- Pros:
 - No need to **touch or trust** end systems
 - Can “bolt on” security
 - **Cheap**: cover many systems w/ single monitor
 - **Cheap**: centralized management

Network-Based Detection

- Issues:
 - Scan for “/etc/passwd”?
 - What about *other* sensitive files?
 - Scan for “.../..”?
 - Sometimes seen in legit. requests (= *false positive*)
 - What about “%2e%2e%2f%2e%2e%2f”? (= *evasion*)
 - It needs to do full HTTP parsing
 - What about “..///.///..///”?
 - It needs to understand Unix filename semantics too!
 - What if it’s HTTPS and not HTTP?
 - Need access to decrypted text / session key – yuck!

Structure of FooCorp Web Services



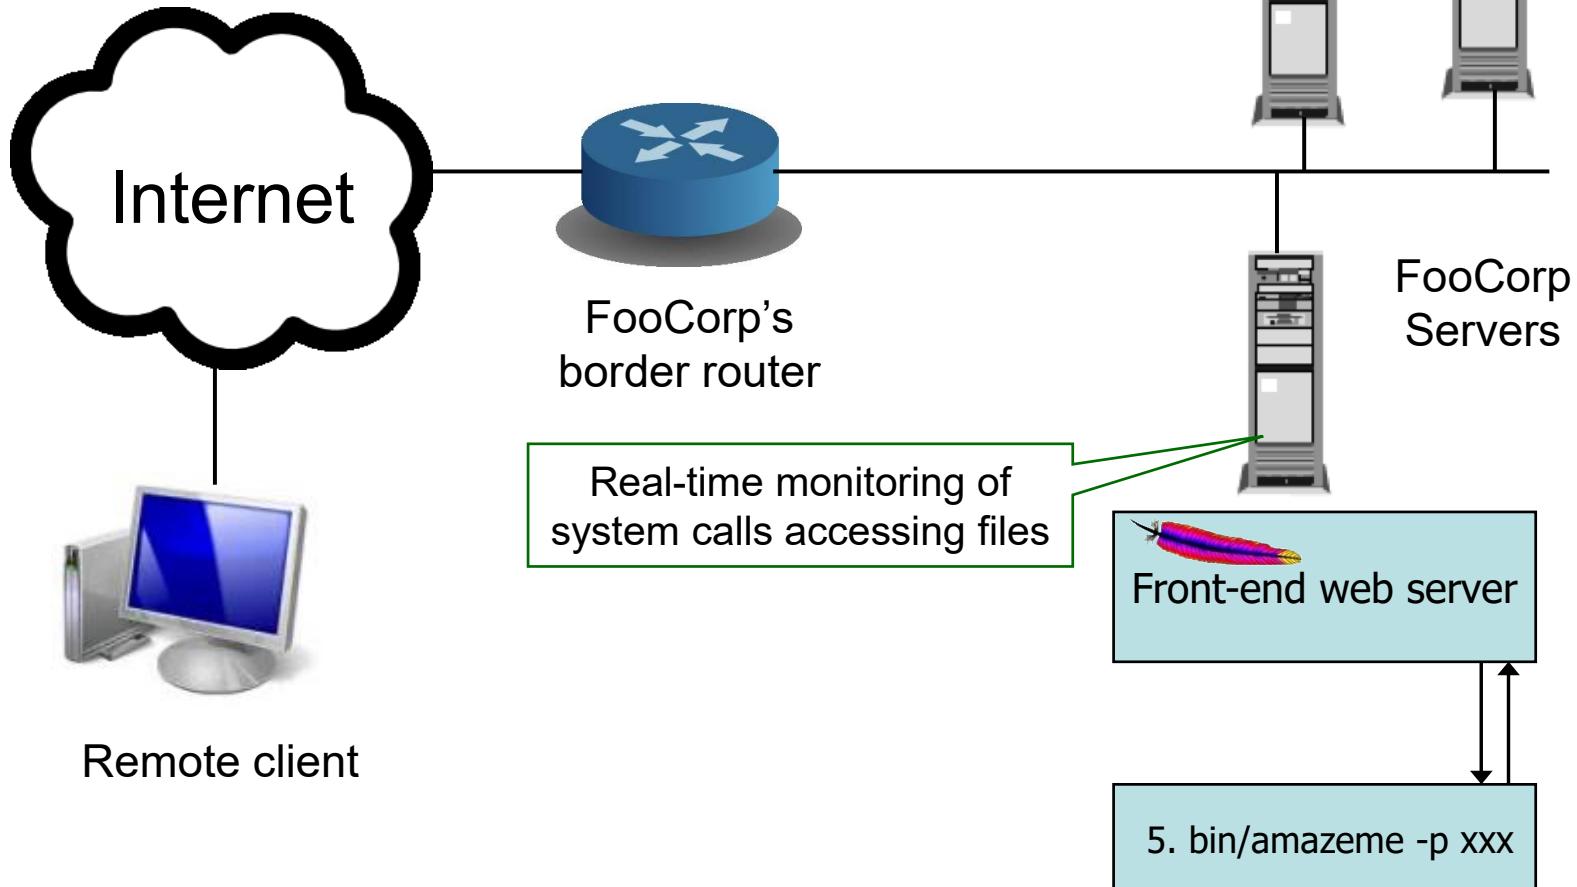
Host-based Intrusion Detection

- Approach #2: instrument the web server
 - Host-based IDS (sometimes called “HIDS”)
 - Resides on a single system and monitors activity on that machine (e.g., OS calls, system logs) and monitors abnormal activity
 - Use heuristics for what is considered to be abnormal activity, e.g., accessing system logs
 - Scan arguments sent to back-end programs
 - Look for “/etc/passwd” and/or “..../”

HIDS

- HIDS attempt #1: scan for arguments sent to back-end programs
 - Look for “/etc/passwd” and/or “..../”
- Pros:
 - No problems with HTTP complexities like %-escapes
 - Works for encrypted HTTPS! (because it gets decrypted at endpoint host)
- Issues:
 - Have to add code to each (possibly different) web server
 - And that effort only helps with detecting web server attacks
 - Still have to consider Unix filename semantics (“..///.//”)
 - Still have to consider other sensitive files

Structure of FooCorp Web Services



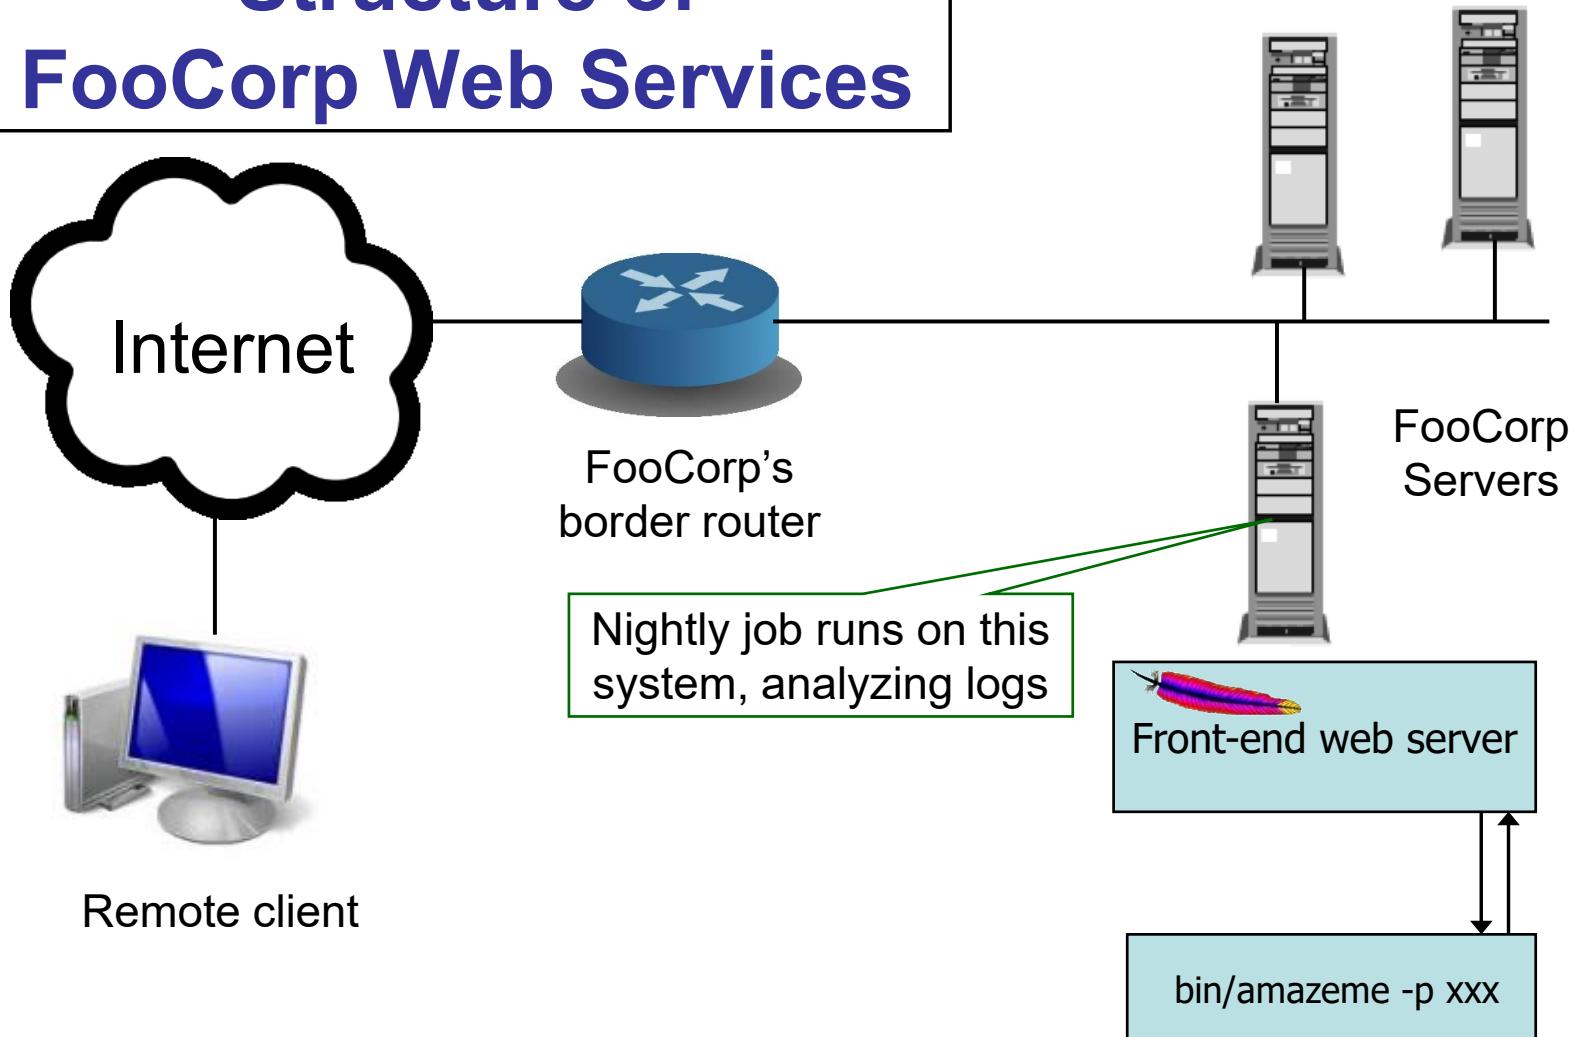
Add system Call Monitoring to HIDS

- HIDS attempt #2: monitor **system call activity** of backend processes
 - Look for access to /etc/passwd which is a sys call

System Call Monitoring (HIDS)

- Approach #2: monitor system call activity of backend processes
 - Look for access to /etc/passwd
- Pros:
 - No issues with any HTTP complexities
 - May avoid issues with filename tricks
 - Attack only leads to an “**alert**” if attack succeeded
 - Sensitive file was indeed accessed
- Issues:
 - Maybe other processes make **legit** accesses to the sensitive files (**false positives**)
 - Maybe we’d like to detect attempts even if they fail?
 - “situational awareness”

Structure of FooCorp Web Services



Log Analysis

- HIDS attempt #3: each night, script runs to analyze **log files** generated by web servers
 - Again scan arguments sent to back-end programs

Log Analysis

- HIDS attempt #3: each night, script runs to analyze log files generated by web servers
 - Again scan arguments sent to back-end programs
- Pros:
 - Cheap: web servers generally already have such logging facilities built into them
 - No problems like %-escapes, encrypted HTTPS since it is at the web application level
- Issues:
 - Again must consider filename tricks, other sensitive files
 - Can't block attacks & prevent from happening
 - Detection delayed, so attack damage may compound
 - If the attack is a compromise, then malware might be able to alter the logs before they're analyzed
 - (Not a problem for directory traversal information leak example)

Typical HIDS

- A combination of the three attempts, monitor system calls, program inputs and system logs. The more information the better.

Detection Accuracy

- Two types of detector errors:
 - **False positive** (FP): alerting about a problem when in fact there was no problem
 - **False negative** (FN): failing to alert about a problem when in fact there was a problem
- Detector accuracy is often assessed in terms of rates at which these occur:
 - Define \square to be the event of an instance of intrusive behavior occurring (something we want to detect)
 - Define \blacksquare to be the event of detector generating alarm
- Define:
 - *False positive rate* = $P[\square | \neg \blacksquare]$
 - *False negative rate* = $P[\neg \square | \blacksquare]$

Perfect Detection

- Is it possible to build a detector for our example with a false negative rate of **0%**?

- Algorithm to detect bad URLs with **0% FN rate**:

```
void my_detector_that_never_misses(char *URL)
{
    printf("yep, it's an attack!\n");
}
```

- In fact, it works for detecting **any** bad activity with no false negatives! **Woo-hoo!**

- Wow, so what about a detector for bad URLs that has **NO FALSE POSITIVES**?!

- `printf("nope, not an attack\n");`

Detection Tradeoffs

- The art of a good detector is achieving an **effective balance** between FPs and FNs
- Suppose our detector has an FP rate of 0.1% and an FN rate of 2%. Is it good enough? Which is better, a very low FP rate or a very low FN rate?
 - Depends on the **cost** of each type of error ...
 - E.g., FP might lead to paging a duty officer and consuming hour of their time; FN might lead to \$10K cleaning up compromised system that was missed
 - ... but also **critically** depends on the rate at which actual attacks occur in your environment

Base Rate Fallacy

- Suppose our detector has a FP rate of 0.1% (!) and a FN rate of 2% (not bad!)
- Scenario #1: our server receives 1,000 URLs/day, and 5 of them are attacks
 -      
 - Expected # FNs each day = $2\% * 5 = 0.1$ (< 1/week)
 - Pretty good!
- Scenario #2: our server receives 10,000,000 URLs/day, and 5 of them are attacks
 -      
- *Nothing changed about the detector; only our environment changed*
 - Accurate detection very challenging when **base rate** of activity we want to detect is quite low

Styles of Detection: Signature-Based

- Idea: look for activity that matches the structure of a **known attack**
- Example (from the freeware *Snort* NIDS):

```
alert tcp $EXTERNAL_NET any -> $HOME_NET  
    139 flow:to_server,established  
    content:"|eb2f 5feb 4a5e 89fb 893e 89f2|"  
    msg:"EXPLOIT x86 linux samba overflow"  
    reference:bugtraq,1816  
    reference:cve,CVE-1999-0811  
    classtype:attempted-admin
```

- Can be at **different semantic layers**
e.g.: IP/TCP header fields; packet payload; URLs

Signature-Based Detection

- E.g. for FooCorp, search for “`..../..`” or “`/etc/passwd`”
- What’s nice about this approach?
 - Conceptually **simple**
 - Takes care of known attacks (of which there are **zillions**)
 - Easy to **share** signatures, build up libraries
- What’s problematic about this approach?
 - Blind to **novel attacks**
 - Might even miss *variants* of known attacks (“`..///.//../`”)
 - Of which there are **zillions**
 - Simpler versions look at low-level **syntax**, not **semantics**
 - Can lead to weak power (either misses variants, or generates lots of **false positives**)

Vulnerability Signatures

- Idea: don't match on known attacks, match on **known problems**
- Example (also from *Snort*):

```
alert tcp $EXTERNAL_NET any -> $HTTP_SERVERS 80
  uricontent: ".ida?"; nocase; dsize: > 239;
  msg:"Web-IIS ISAPI .ida attempt"
  reference:bugtraq,1816
  reference:cve,CAN-2000-0071
  classtype:attempted-admin
```
- That is, match URLs that invoke ***.ida?***, have more than **239 bytes** of payload
- This example detects any* attempt to exploit a particular buffer overflow in IIS web servers
 - Used by the “Code Red” worm
- * (Note, signature is not quite complete)

Vulnerability Signatures

- What's nice about this approach?
 - Conceptually fairly simple *Benefits of attack signatures*
 - Takes care of known attacks
 - Easy to share signatures, build up libraries
 - Can detect **variants** of known attacks
 - Much more **concise** than per-attack signatures
- What's problematic?
 - Can't detect **novel attacks** (new vulnerabilities)
 - Signatures can be **hard** to write / express
 - Can't just observe an attack that works ...
 - ... need to delve into **how** it works

Styles of Detection: Anomaly-Based

- Idea: attacks look **peculiar**.
- High-level approach: develop a **model** of **normal** behavior (say based on analyzing historical logs). Flag activity that **deviates** from it.
- FooCorp example: maybe look at distribution of characters in URL parameters, learn that some are rare and/or don't occur repeatedly
 - If we happen to learn that '.'s have this property, then could detect the attack **even without knowing it exists**
- Big benefit: potential detection of a wide range of attacks, **including novel ones**

Anomaly Detection

- What's problematic about this approach?
 - Can **fail to detect** known attacks
 - Can **fail to detect** novel attacks, if don't happen to look peculiar along measured dimension
 - What happens if the historical data you train on includes attacks?
 - **Base Rate Fallacy** particularly acute: *if prevalence of attacks is low, then you're more often going to see benign outliers*
 - **High FP rate**
 - OR: require such a stringent deviation from “normal” that most attacks are missed (**high FN rate**)

Hard to make work well - not widely used today

Specification-Based Detection

- Idea: don't learn what's normal; specify what's allowed
- FooCorp example: decide that all URL parameters sent to foocorp.com servers must have at most one '/' in them
 - Flag any arriving param with > 1 slash as an attack
- What's nice about this approach?
 - Can detect novel attacks
 - Can have low false positives
 - If FooCorp audits its web pages to make sure they comply
- What's problematic about this approach?
 - Expensive: lots of labor to derive specifications
 - And keep them up to date as things change ("churn")

Styles of Detection: Behavioral

- Idea: don't look for attacks, look for **evidence of compromise**
- FooCorp example: inspect all output web traffic for any lines that match a `passwd` file
- Example for monitoring user shell keystrokes:
`unset HISTFILE` (don't save bash history)
- Example for catching **code injection**: look at sequences of system calls, flag any that prior analysis of a given program shows it can't generate
 - E.g., observe process executing `read()`, `open()`, `write()`, `fork()`, `exec()` ...
 - ... but there's *no code path* in the (original) program that calls those in exactly that order!

Behavioral-Based Detection

- What's nice about this approach?
 - Can detect a wide range of **novel** attacks
 - Can have **low false positives**
 - Depending on degree to which behavior is distinctive
 - E.g., for system call profiling: **no false positives!**
 - Can be **cheap** to implement
 - E.g., system call profiling can be mechanized
- What's problematic about this approach?
 - Post facto detection: discovers that you definitely have a problem, w/ **no opportunity to prevent it**
 - **Brittle**: for some behaviors, attacker can maybe avoid it
 - Easy enough to not type “unset HISTFILE”
 - How could they evade system call profiling?
 - **Mimicry**: adapt injected code to comply w/ allowed call sequences

The Problem of Evasion

- For any detection approach, we need to consider how an adversary might (try to) **elude** it
 - *Note: even if the approach is evadable, it can still be useful to operate in practice*
 - **But:** if it's very easy to evade, that's especially worrisome (security by obscurity)

The Problem of Evasion

- Imperfect observability is particularly acute for network monitoring
- Consider detecting occurrences of the (arbitrary) string “**root**” inside a network connection ...
 - We get a copy of each packet, how hard can it be?

Detecting “root”: Attempt #1

- Method: scan each packet for ‘r’, ‘o’, ‘o’, ‘t’
 - Perhaps using Boyer-Moore, Aho-Corasick, Bloom filters ...

Packet

1**root**.....

Are we done?

Oops: TCP *doesn't preserve text boundaries*

1**ro**

2 **ot**.....

Packet #1

Packet #2

Fix?

Detecting “root”: Attempt #2

- Okay: remember match from end of previous packet

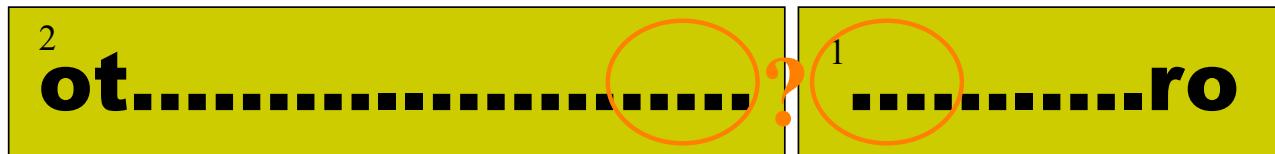


When 2nd packet arrives, continue working on the match

- Now we're managing **state** :-(

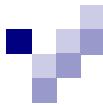
Are we done?

Oops: IP doesn't guarantee in-order arrival

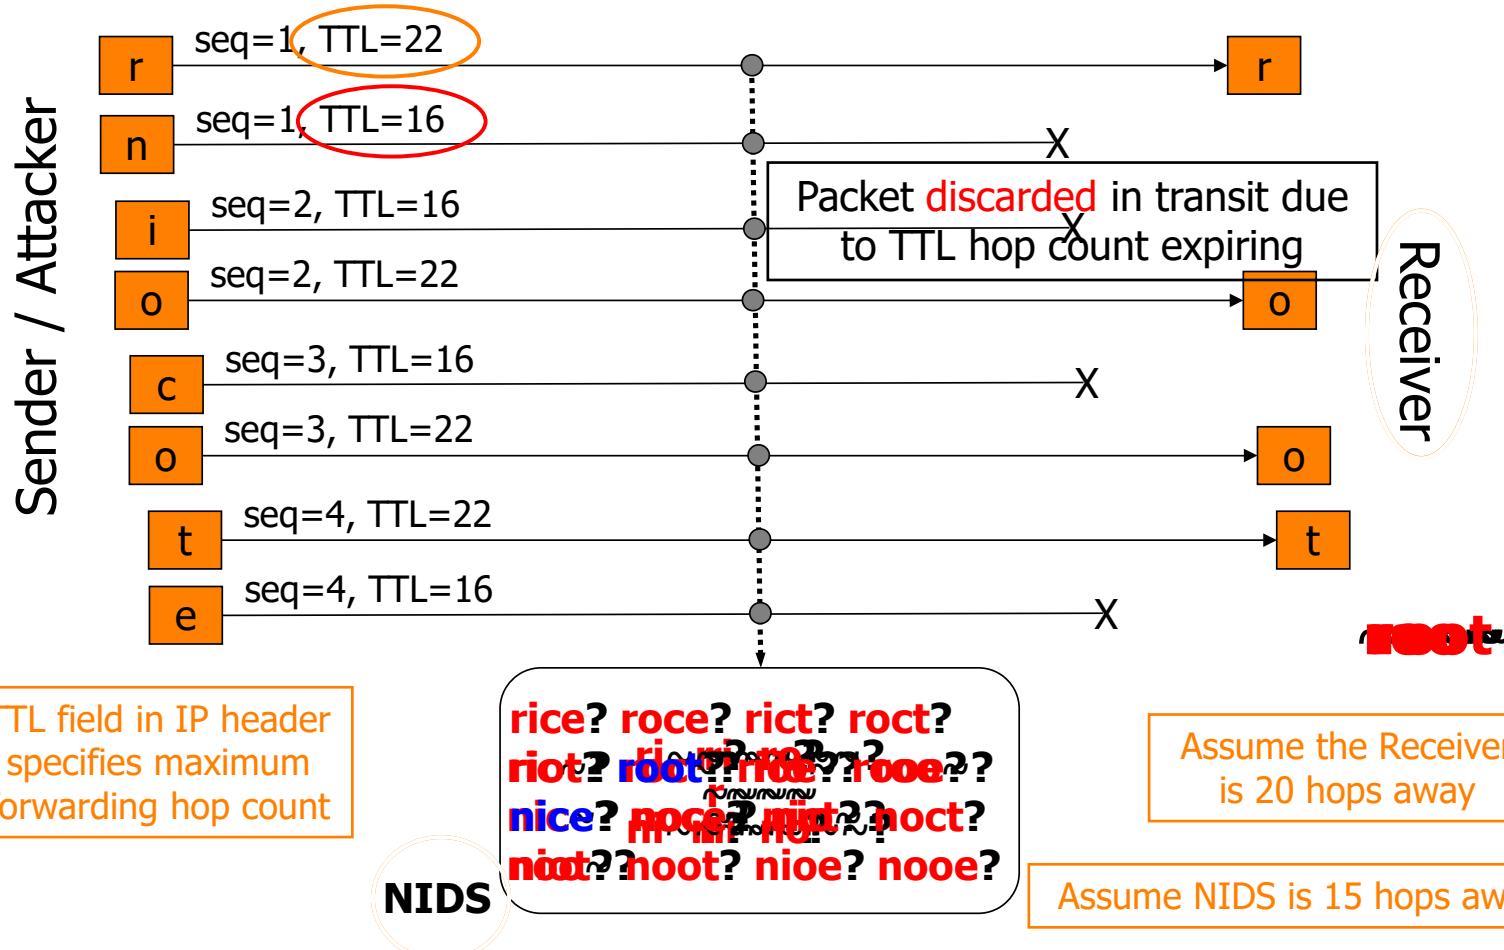


Detecting “root”: Attempt #3

- Fix?
- We need to reassemble the **entire** TCP bytestream
 - Match sequence numbers
 - Buffer packets with later data (above a sequence “hole”)
- Issues?
 - Potentially requires a lot of **state**
 - Plus: attacker can cause us to **exhaust state** by sending lots of data above a sequence hole
- But at least we’re done, right?



Full TCP Reassembly is Not Enough



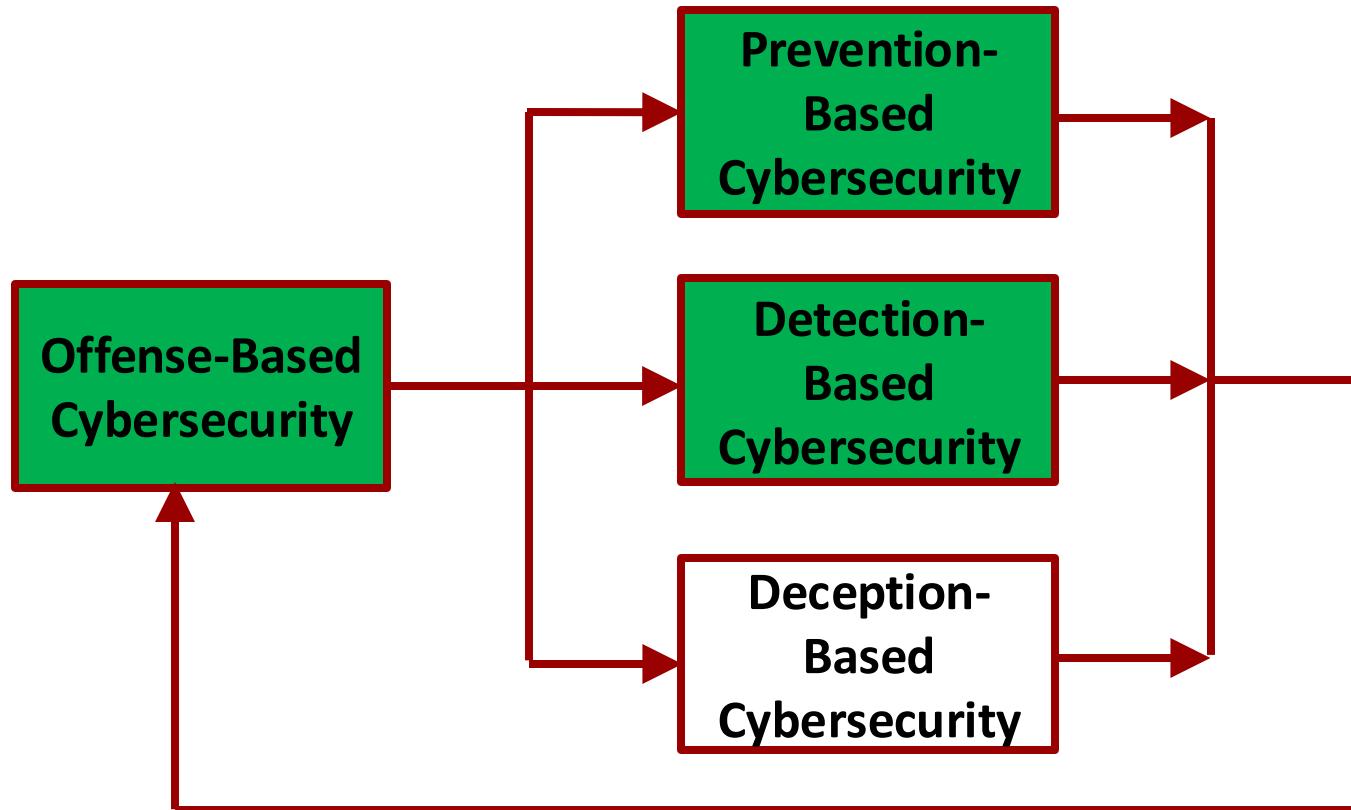
Inconsistent TCP Retransmissions

- Fix?
- Idea: NIDS can **alert** upon seeing a retransmission inconsistency (two packets for same seqno), as surely it reflects someone up to no good
- This **doesn't work well in practice**: TCP retransmissions broken in this fashion occur in live traffic
 - Fairly rare (23 times in a day of ICSI traffic)
 - But real evasions **much rarer still** (Base Rate Fallacy)
 - This is a *general problem* with alerting on such ambiguities
- Idea: if NIDS sees such a connection, **kill it**
 - Works for this case, since benign instance is already fatally broken
 - But for other evasions, such actions have **collateral damage**
- Idea: **rewrite** traffic to remove ambiguities
 - Works for network- & transport-layer ambiguities
 - But must operate **in-line** and **at line speed**

Summary of Evasion Issues

- Evasions arise from **uncertainty** (or **incompleteness**) because detector must infer behavior/processing it can't directly observe
 - A general problem any time detection separate from potential target
- One general strategy: impose canonical form ("normalize")
 - E.g., rewrite URLs to expand/remove hex escapes
 - E.g., enforce blog comments to only have certain HTML tags
- (Another strategy: analyze **all** possible interpretations rather than assuming one)
 - E.g., analyze raw URL, hex-escaped URL, doubly-escaped URL ...)
- Another strategy: fix the basic observation problem
 - E.g., monitor **directly** at end systems

Course structure



End of Lecture 17