



Project 2: **25** days left

Deception-Based Cyber Security: Honeypot

CS 459/559: Science of Cyber Security
19th Lecture

Instructor:

Guanhua Yan

Agenda

- ~~Quiz 1: September 29 (closed book)~~
- ~~Project 1 (offense): October 10~~
- Course wrapup & presentation lottery:
November 10
- Quiz 2: November 12
- Presentations: 11/17, 11/19, 11/24, 12/1,
12/3
- CTF competition: November 26
- Project 2 (defense): December 5
- Final report: December 15



CTF leaderboard

	User Name	Successful Attack	Score
1	sandworm	Buffer Overflow, DNS Tunneling, Known Plaintext Attack, Network Reconnaissance, Program Wrapper, Reverse Proxy, SQL Injection, Tiny Shell Exploit	120
2	slee	Buffer Overflow, DNS Tunneling, Network Reconnaissance, Program Wrapper, Reverse Proxy, SQL Injection, Tiny Shell Exploit	110
3	jeff	Buffer Overflow, DNS Tunneling, Network Reconnaissance, Program Wrapper, Reverse Proxy, SQL Injection, Tiny Shell Exploit	110
4	Sandeep	DNS Tunneling, Network Reconnaissance, Reverse Proxy, SQL Injection	80
5	Srimunagala	DNS Tunneling, Network Reconnaissance, Reverse Proxy, SQL Injection	80
6	akoval	Network Reconnaissance, Reverse Proxy, SQL Injection	70
7	haritha	DNS Tunneling, Network Reconnaissance, Reverse Proxy	70
8	JamesRatanDukkipati	DNS Tunneling, Program Wrapper(partial), SQL Injection	70

Outline

- **What is honeypot?**
- **Honeypot types**
- **Honeypot uses**
- **Honeypot challenges**
- **Honeypot examples**
 - Honeyd
 - Honeynets

What is honeypot?

Motivation

- Key to effective intrusion detection is **information**
 - Learn more about past attacks
 - Detect currently occurring attacks
 - Identify new types of attacks
 - Do all this in real time

What is a honeypot?

“A honeypot is an information system resource whose value lies in unauthorized or illicit use of that resource.”

-- Lance Spitzer
(Founder, Honeynet Project)

- Could be...
 - A password file
 - An Excel spreadsheet
 - An entry in a database
 - A computer on a network ← This is the kind of honeypot we will talk about!

What is a honeypot?

- The basic idea: set up a “normal” but unused computer on your network
 - Nobody knows it’s there, so it should get no legitimate network traffic
 - Any traffic it gets is malicious by definition
 - All interactions with the honeypot are logged on a remote machine

What is a honeypot?

- Advantages of using a honeypot
 - Small, valuable data sets: no normal traffic, only attacks
 - Very few false positives or false negatives
 - Uses minimal resources
 - Easy to set up and use
 - Can capture new types of attacks
 - Can gather detailed information about attacks

Honeypot types

Types of honeypots

- To an attacker, a honeypot should always look like a normal computer – but what is it really?
 - It could actually be a normal computer
 - It could be a simulation of certain aspects of a computer
 - Different types of honeypots are useful for different purposes

Types of honeypots

- Two basic categories:
 - Low-interaction honeypots
 - High-interaction honeypots

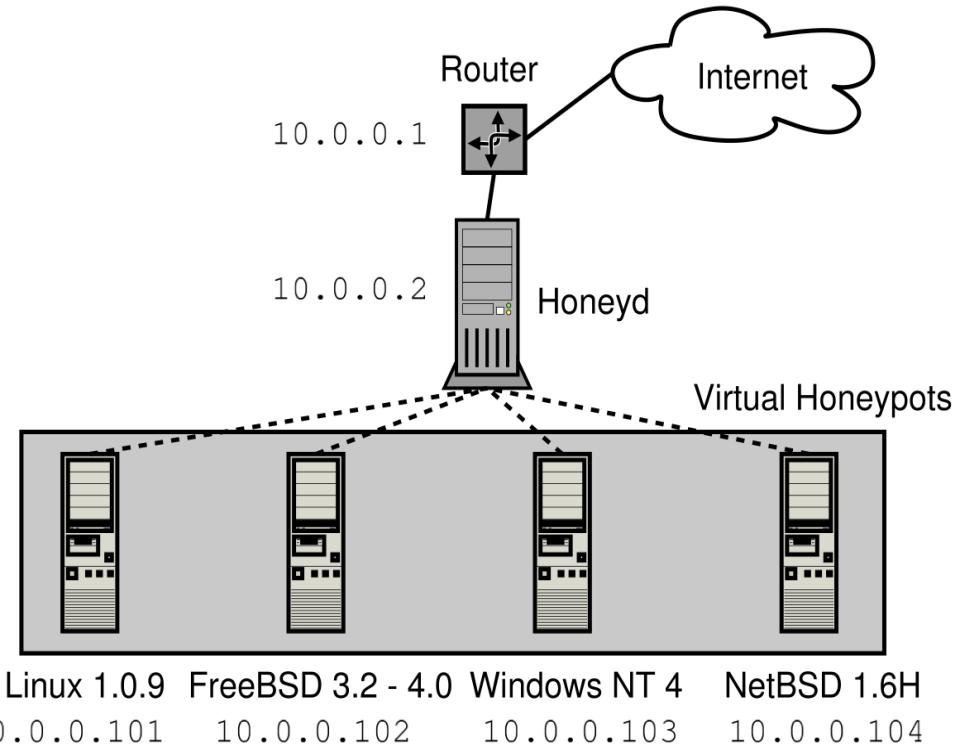
Low-interaction honeypots

- Attacker interacts with a “simulated” computer
- Many levels of simulation possible
 - Network stack
 - Services
 - Operating system

Low-interaction honeypots

- One real machine can simulate a whole network of virtual honeypots

Architecture of Honeyd,
a low-interaction honeypot.
Only the router and the Honeyd
machine (10.0.0.2) are real
computers!



Low-interaction honeypots

- Advantages
 - Very simple
 - Low-risk (attacker never gets into a real system)
 - Require very minimal resources
- Disadvantages
 - Only collect limited information
 - Might not detect new types of attacks
 - Easy for attacker to detect

High-interaction honeypots

- Real machines running real services
- We assume that these machines *will* be compromised!
 - All interactions with the machines are monitored and logged, providing detailed information about what the attacker did

High-interaction honeypots

- Fishbowl analogy

- Set up a framework that provides data logging and security (the fishbowl)
- Within that framework, put machines that you want the attacker to interact with (the rocks, plants, etc)
- Watch how the attacker (the fish) interacts with the machines



High-interaction honeypots

- Two main requirements of this framework
 - *Data Control* – prevent the attacker from using the honeypots to harm other machines
 - *Data Capture* – record all the attacker's activities
- Both of these should be invisible to the attacker!

High-interaction honeypots

- Advantages
 - Capture a detailed profile of an attack
 - Can capture new types of attacks
- Disadvantages
 - Difficult to set up a good high-interaction honeypot
 - May put other machines in your network at risk
 - Monitoring the honeypots is time-intensive

Honeypot uses

Uses of honeypots

- What can you do with a honeypot?
- Intrusion detection/prevention
 - Lots of ways to use a honeypot as part of your security system
 - Most honeypot research is in this area
- Attack analysis
 - Observe attackers' behavior and develop better tools to guard against it

Uses of honeypots

- Decoys
 - Populate all unused addresses on your network with honeypots
 - Attacker has to waste time trying to attack the honeypots
 - Slows down the spread of worms
 - Slows down and annoys human attackers (maybe enough to make them go away?)

Uses of honeypots

- Tarpits
 - Intended to slow an attacker down
 - Labrea Tarpit
 - Allows attacker to open a TCP connection, then reduces window size to 0
 - Attacker can't get any data through, and can't close the connection
 - Connection uses up resources on the attacker's system

Uses of honeypots

- Tarpits (continued)
 - Open mail relays
 - The honeypot offers an anonymous mail relay (which attracts spammers)
 - Responds very slowly to SMTP commands
 - Forces spammers to waste time interacting with the honeypot
 - Honeypot may pretend to forward the mail, but actually drop it

Uses of honeypots

- Burglar alarms
 - When the honeypot is compromised, admins know that an attack is going on in their network
 - Honeypot logs provide detailed information about the attack
 - Some evidence (from GT Honeynet) that attacks can be predicted a few days in advance, based on abnormal activity on the honeypots

Uses of honeypots



- Automatic signature generation
 - Honeycomb – a plug-in for honeyd
 - Detects patterns in the logged data, creates Snort and Bro signatures
 - Works fairly well with no human input, and much faster than manual signature generation

Uses of honeypots

- Many more ways to use honeypots
 - Identify zero-day worms
 - Disrupt DDoS attacks
 - Monitor botnets
 - Etc...

Honeypot challenges

Problems with honeypots

- So what's wrong with honeypots?
 - Attacker may do bad things with the compromised system
 - Attacker may discover that the system is a honeypot
 - Legal concerns
 - Difficult to catch more intelligent attackers with honeypots

Problems with honeypots

- Once a honeypot is compromised...
 - It may be used to attack other machines (on your network or elsewhere).
 - Preventing this should be the top priority of a honeynet – but no guarantees!
 - It may be used for criminal activity (ex. serving illegal files)
 - If any of this is detected, it will initially be blamed on you!

Problems with honeypots

- Legal concerns
 - Privacy – anybody interacting with the honeypot does not know that the interactions are being logged

Problems with honeypots

- Legal concerns
 - Liability – if your honeypot is used to attack someone else, can they sue you?
 - You intentionally allowed the attacker to get in, so you may be blamed

Problems with honeypots

- What kind of attackers can a honeypot catch?
 - It depends on the “bait” you use
 - Normal machines will mostly attract automated attacks
 - To catch specific threats (like credit card thieves) you need a honeypot that “looks” valuable to them!
 - This is very hard to do, so it’s hardly ever done!



Honeypot examples

Examples of honeypots

- Honeyd (a low-interaction honeypot)
- Honeynets (a high-interaction honeypot)

Honeyd

- Low-interaction honeypot
- Runs on a single computer
 - Simulates a group of virtual machines
 - Simulates the physical network between them
- Simulates only the network stack of each machine
- Intended primarily to fool fingerprinting tools

Honeyd

- Fingerprinting
 - Attackers often try to learn more about a system before attacking it
 - Can determine a machine's operating system by “testing” its network behavior
 - How the initial TCP sequence number is created
 - Response packets for open and closed ports
 - Configuration of packet headers
 - Common fingerprinting tools: Xprobe, Nmap

Honeyd

- Setting up Honeyd
 - Configure the Honeyd machine to receive packets addressed to the virtual machines
 - Several ways to do this:
 - Add routes in routing table
 - Proxy ARP
 - Network tunneling

Honeyd

- Honeyd logs all received packets
- For TCP, UDP, and ICMP packets:
 - Sends an appropriate response packet
 - Adjusts the packet content so it looks like it came from the virtual machine
- This response is determined by the config file!

Honeyd

- A Honeyd config file:

```
create windows
```

Define the OS (this refers to an nmap fingerprint!)

```
set windows personality "Windows NT 4.0 Server SP5-SP6"
```

```
set windows default tcp action reset
```

How to respond to incoming packets

```
set windows default udp action reset
```

```
add windows tcp port 80 "perl scripts/iis-0.95/iisemul8.pl"
```

Run a script to emulate a web server

```
add windows tcp port 139 open
```

Set open ports

```
add windows tcp port 137 open
```

```
add windows udp port 137 open
```

```
add windows udp port 135 open
```

Set machine's uptime

```
set windows uptime 3284460
```

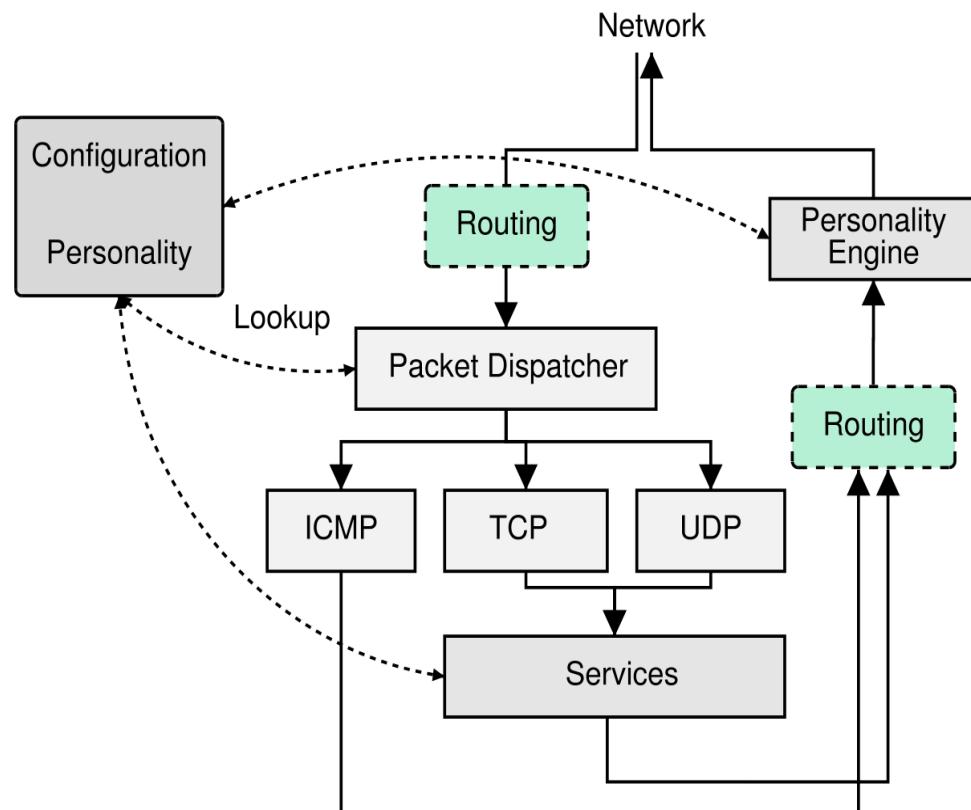
```
bind 192.168.1.201 windows
```

Bind this machine to an IP address

Honeyd

- Honeyd architecture

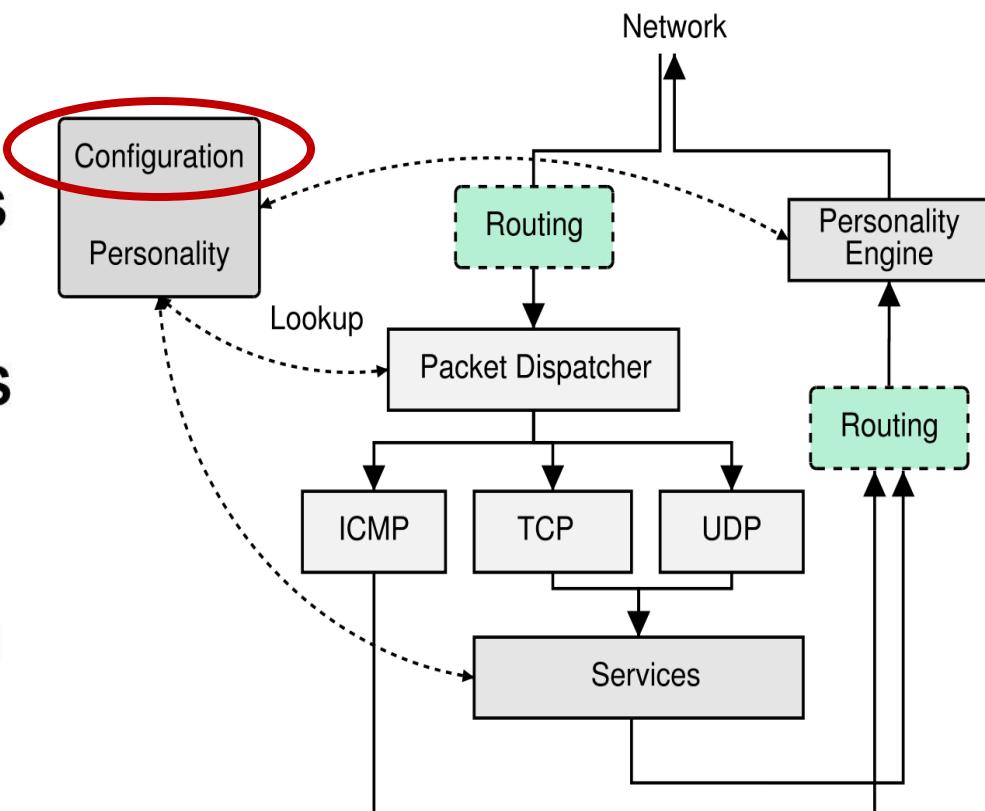
- Packet dispatcher
- Configuration database
- Protocol handlers
- Router (maybe)
- Personality engine



Honeyd

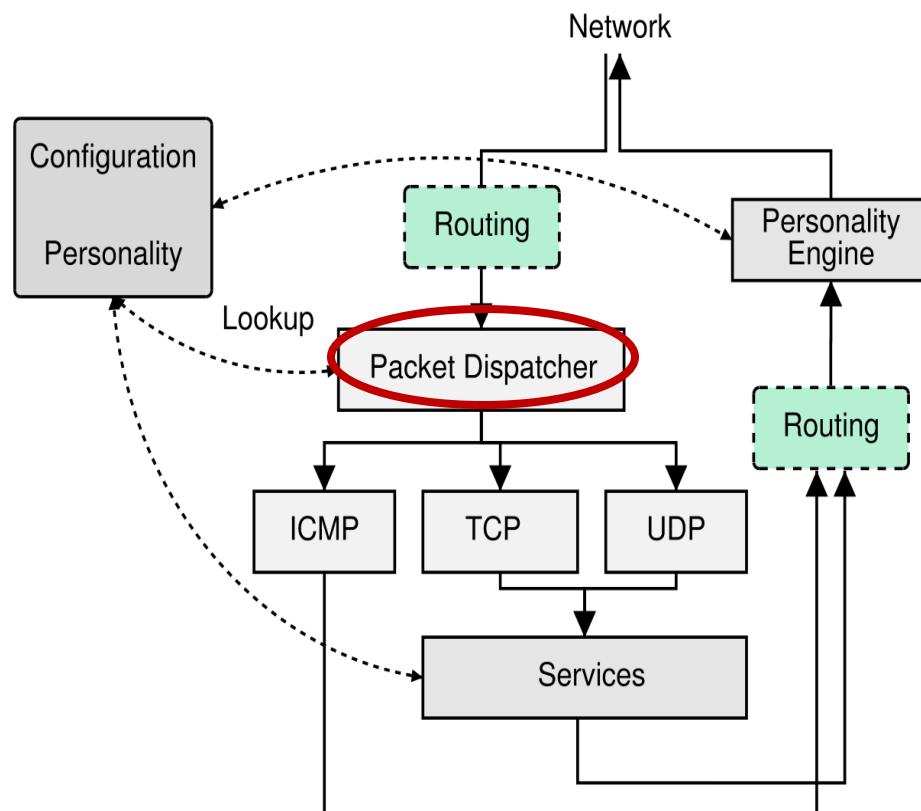
- Configuration database

- A list of configurations like the one we saw
- Links virtual machines to IP addresses
- Uses a default template if no specific configuration is available



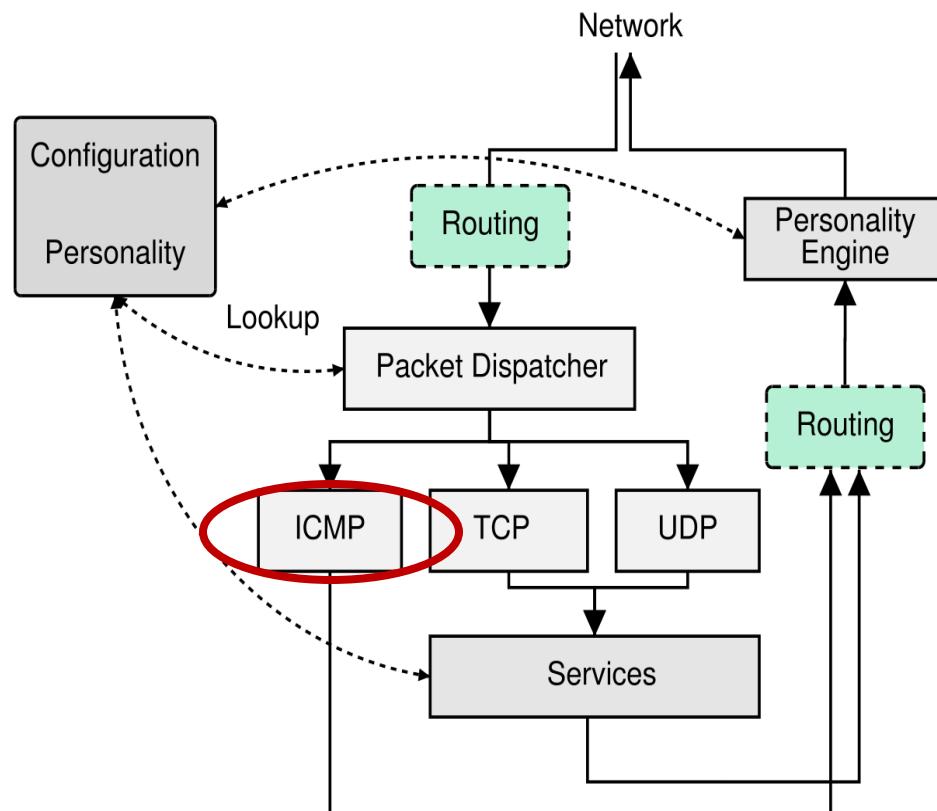
Honeyd

- **Packet dispatcher**
 - Processes incoming packets
 - Looks up the configuration of the virtual machine for each packet
 - Passes TCP, UDP, and ICMP packets to the correct protocol handlers (along with configuration)
 - Drops all packets from other protocols



Honeyd

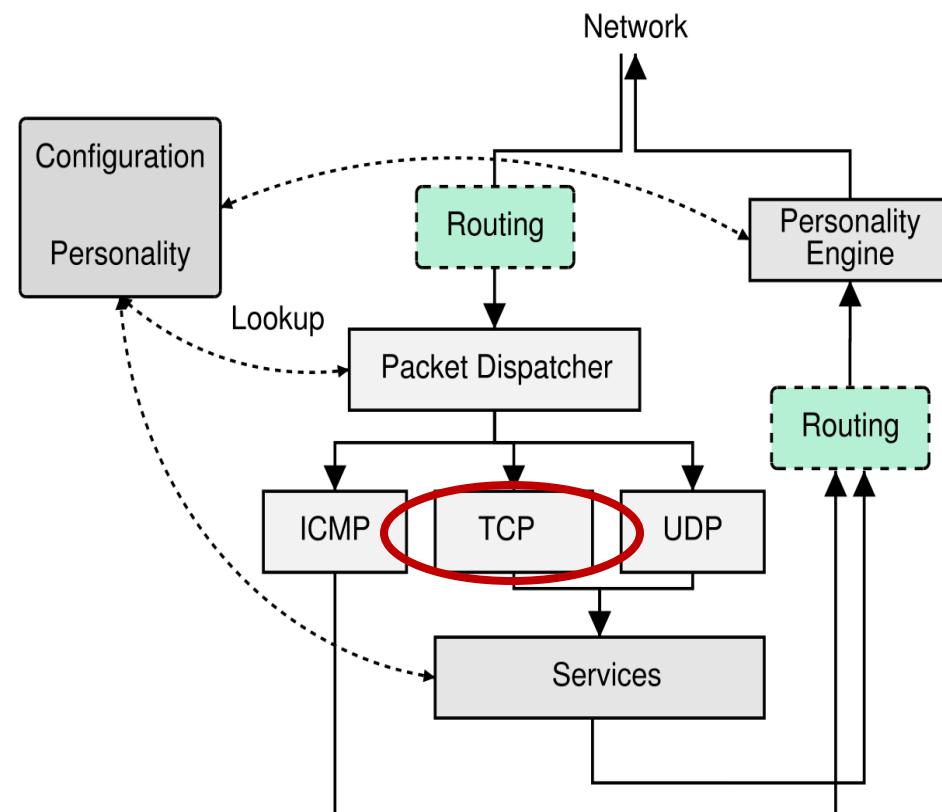
- Protocol handlers:
ICMP
 - Responds to echo requests
 - May respond to other types of requests, depending on configuration



Honeyd

- Protocol handlers:
TCP

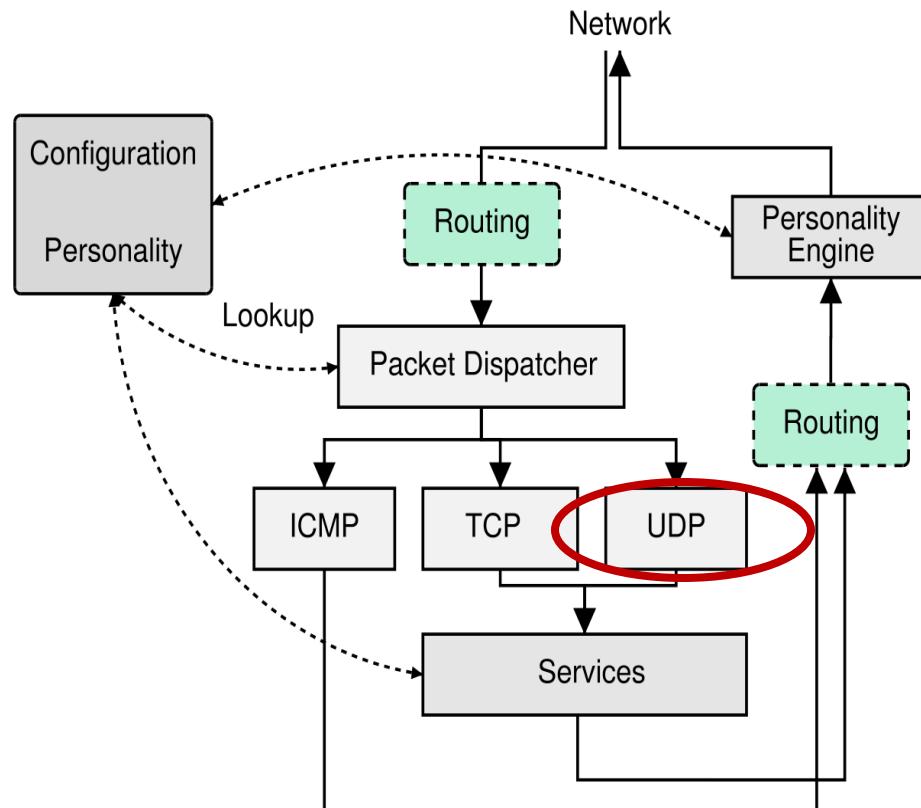
- Implements connection establishment and teardown
- Passes packets to simulated “services”



Honeyd

- Protocol handlers:
UDP

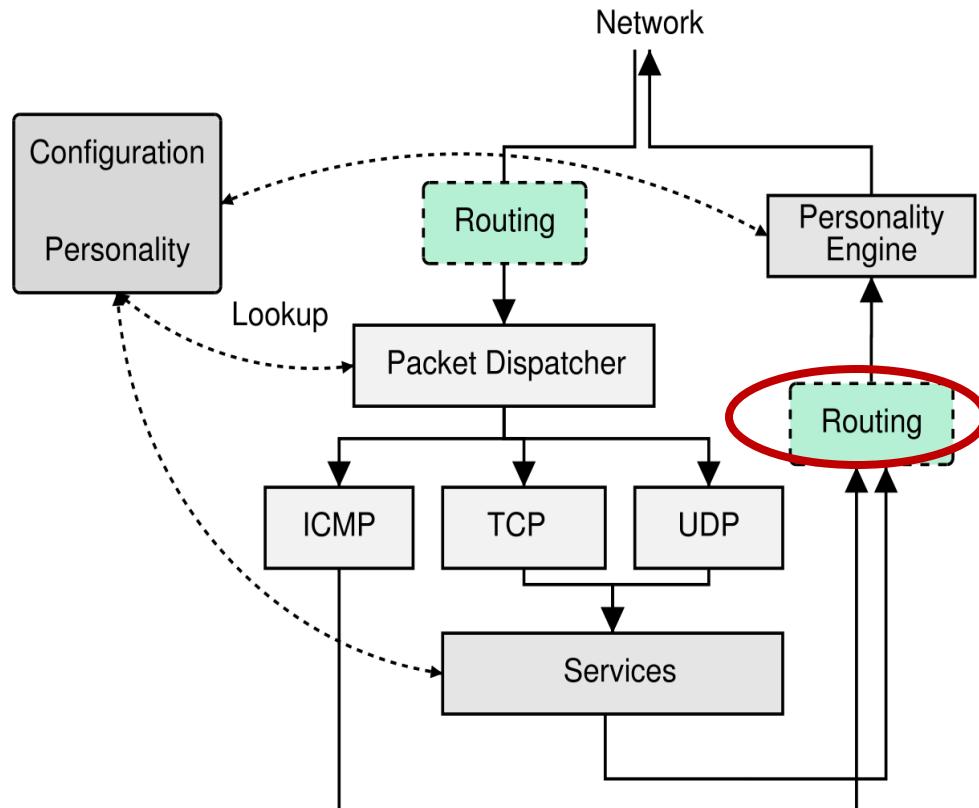
- If port is open,
passes packet to
the appropriate
“service”
- If port is closed,
sends an ICMP *port
unreachable*
message



Honeyd

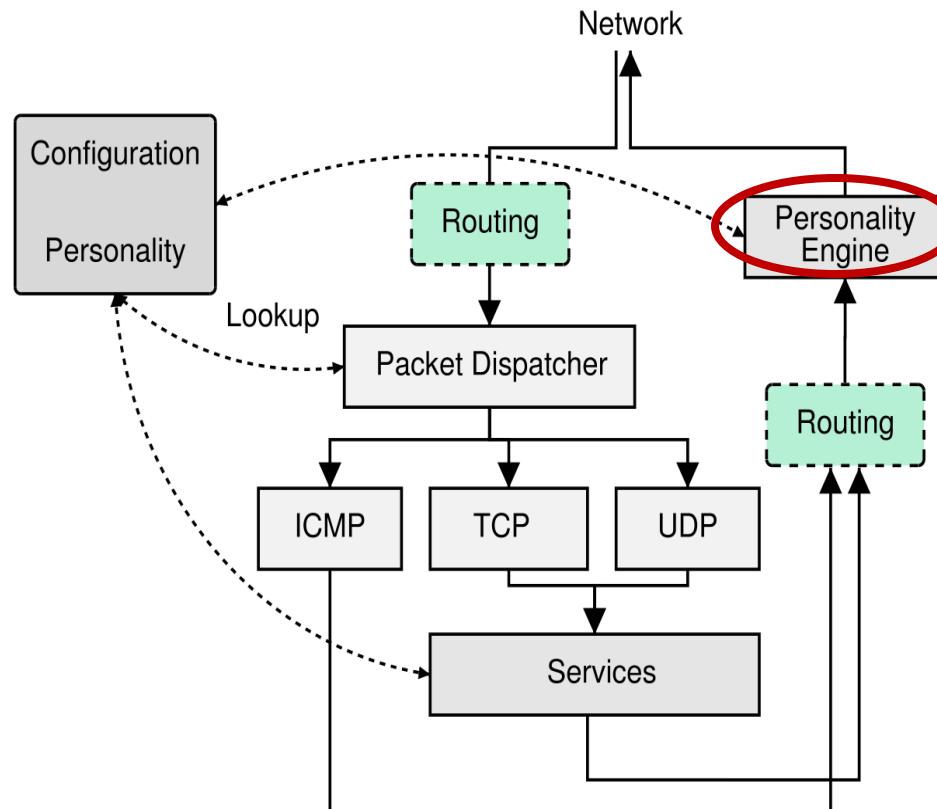
- Router

- Simulates a routing tree between the virtual machines
- Simulates latency and packet loss
- Decrement packet's TTL field



Honeyd

- Personality engine
 - Looks at outgoing packets just before they are sent
 - Adjusts packet headers so that their fingerprints will be correct



Honeyd

- Simulating a routing topology
 - Network links can also be defined in the configuration file
 - Incoming and outgoing packets traverse a virtual routing tree
 - Packets may be delayed or dropped, to simulate latency and loss
 - Real machines can be integrated into this topology

Honeyd

- Logging
 - Honeyd logs all attempted connections on all protocols
 - The “services” should keep their own logs – these usually provide more interesting data
 - All log data is stored on the local machine
 - So it should be secure!

Honeyd

- Uses of Honeyd
 - Network decoy
 - Detecting worms
 - Capturing spam
 - Providing a “front end” that selectively forwards packets to high-interaction honeypots

Honeyd

- Open source, available at
<http://www.honeyd.org>
- Includes sample configurations and scripts
- People have contributed more scripts to simulate different services (and worms)

Honeynets

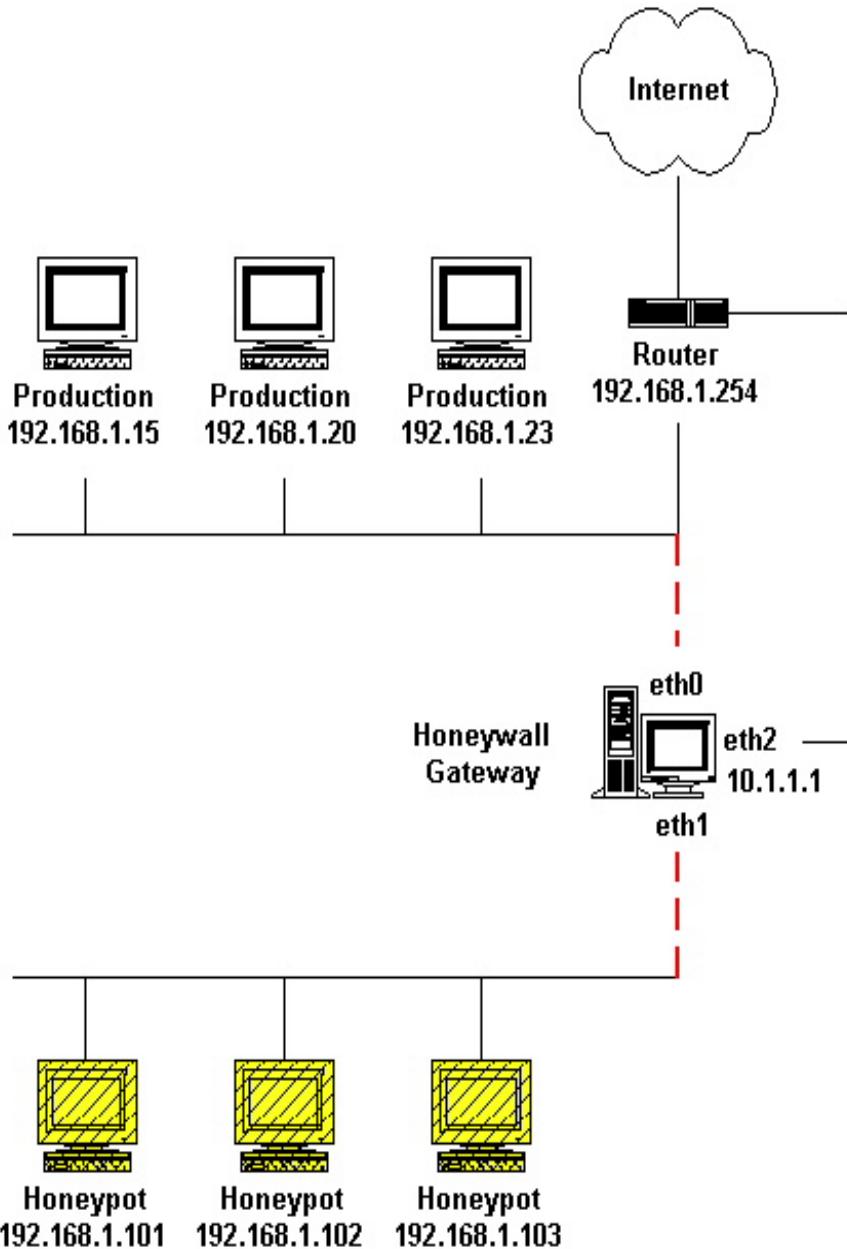
- High-interaction honeypots
- Technically, anything that implements Data Control and Data Capture is a Honeynet
- We will look at a specific architecture: Gen II Honeynets

Honeynets

- Gen II Honeynets
 - Network of real machines (honeypots)
 - Honeywall – a gateway between the honeypots and the rest of the world
 - The Honeywall provides Data Control and Data Capture
 - The Honeywall should be invisible to the attacker!

Honeynets

- Architecture of a GEN II Honeynet
 - Honeywall is a layer 2 bridge
 - eth0: interface to the rest of the network
 - eth1: interface to the Honeynet
 - eth2: for remote administration of Honeywall



Honeynets

- Data Control
 - Goal: prevent an attacker from using a Honeynet system to harm non-Honeynet systems
 - Tradeoff between allowing attackers freedom to act and containing them
 - Should be implemented at several layers
 - Should “fail closed” – if something in the Honeynet architecture fails, the protections should stay up!

Honeynets

- Data Control in Gen II Honeynets
 - Connection counting
 - Traffic scans

Honeynets

- Connection Counting

- Limits the number of outbound connections a honeypot can initiate in a time period
- If connection limit is reached, successive outbound connections are limited over the time scale
- Can be set individually for different protocols

SCALE="day"
TCPRATE="15"
UDPRATE="20"
ICMPRATE="50"
OTERRATE="15"

15 outbound TCP connections are allowed from a single machine; if that limit is hit, 15 more connections will be permitted over the next 24 hours.

Honeynets

- Connection Counting

- Choosing the connection limits is a tradeoff between information and security
- Low connection limit: can be used as a signature to identify the honeypot
- High connection limit: allows attacker to do much more damage!

Honeynets

- Traffic scans
 - Uses snort_inline to scan all packets as they go through the gateway
 - snort_inline: a version of snort that can also drop or modify packets
 - Drops or disables known viruses
 - Custom snort ruleset that focuses on outbound attacks only

Honeynets

- Data Capture
 - Goal: Log all the attacker's activities within the Honeynet, without the attacker noticing
 - Capture information at multiple levels
 - Store all data on a remote, secured system

Honeynets

- Data Capture in Gen II Honeynets
 - Firewall logs
 - Network traffic logs
 - System activity

Honeynets

- Firewall logs
 - Logs all inbound and outbound connections through the Honeywall
 - This is usually the first indication of what an attacker is doing!

Honeynets

- Network traffic logs
 - Logs the complete payload of every packet that goes through the Honeywall
 - Uses a second snort process to do the logging

Honeynets

- System activity
 - Captures the attacker's activity on the honeypot itself
 - Important to log this, since network traffic might be encrypted!
 - Implemented using a kernel patch (Sebek)
 - Logs all system activity
 - Cannot “see” UDP packets with a predefined “magic number”
 - This allows logs to be sent to a remote machine

Honeynets

- Alerting
 - Honeynets are useless if you don't know (preferably right away) when a break-in has occurred
 - Ideally, have a trained admin monitoring the Honeynet at all times
 - Automated monitoring tools – can look for suspicious activity and send out alerts
 - Swatch – a tool that monitors log files for predefined patterns

Honeynets

- The Honeynet Project
 - Developed Gen II Honeynets architecture
 - All tools are open-source and available at <http://www.honeynet.org/>
 - Honeynet Research Alliance: coordinates honeynet research around the world

Summary

- Honeypots gather data about network attacks
- A honeypot has no production value, so all interactions with it are considered attacks
- Honeypots should provide the illusion of being “normal machines” while minimizing risks to the rest of the network

End of Lecture 19